

Lecture 23

Using Nsight Compute and Nsight Systems

Carl Pearson

pearson@illinois.edu

Objective

- Nsight Toolsuite
 - Nsight Systems System-Level Profiling
 - Nsight Compute: Kernel-Level Profiling
- Not discussed
 - Other performance tools: gprof, perf, vtune, ...

Outline

- Introduction to Profiling
- Development Model and Profiling Strategy
- Preparing Application for Profiling
- Measuring time with CUDA Events
- Reminder / Introduction to Matrix Multiplication
- Nvidia Nsight Compute
- Nvidia Nsight Systems

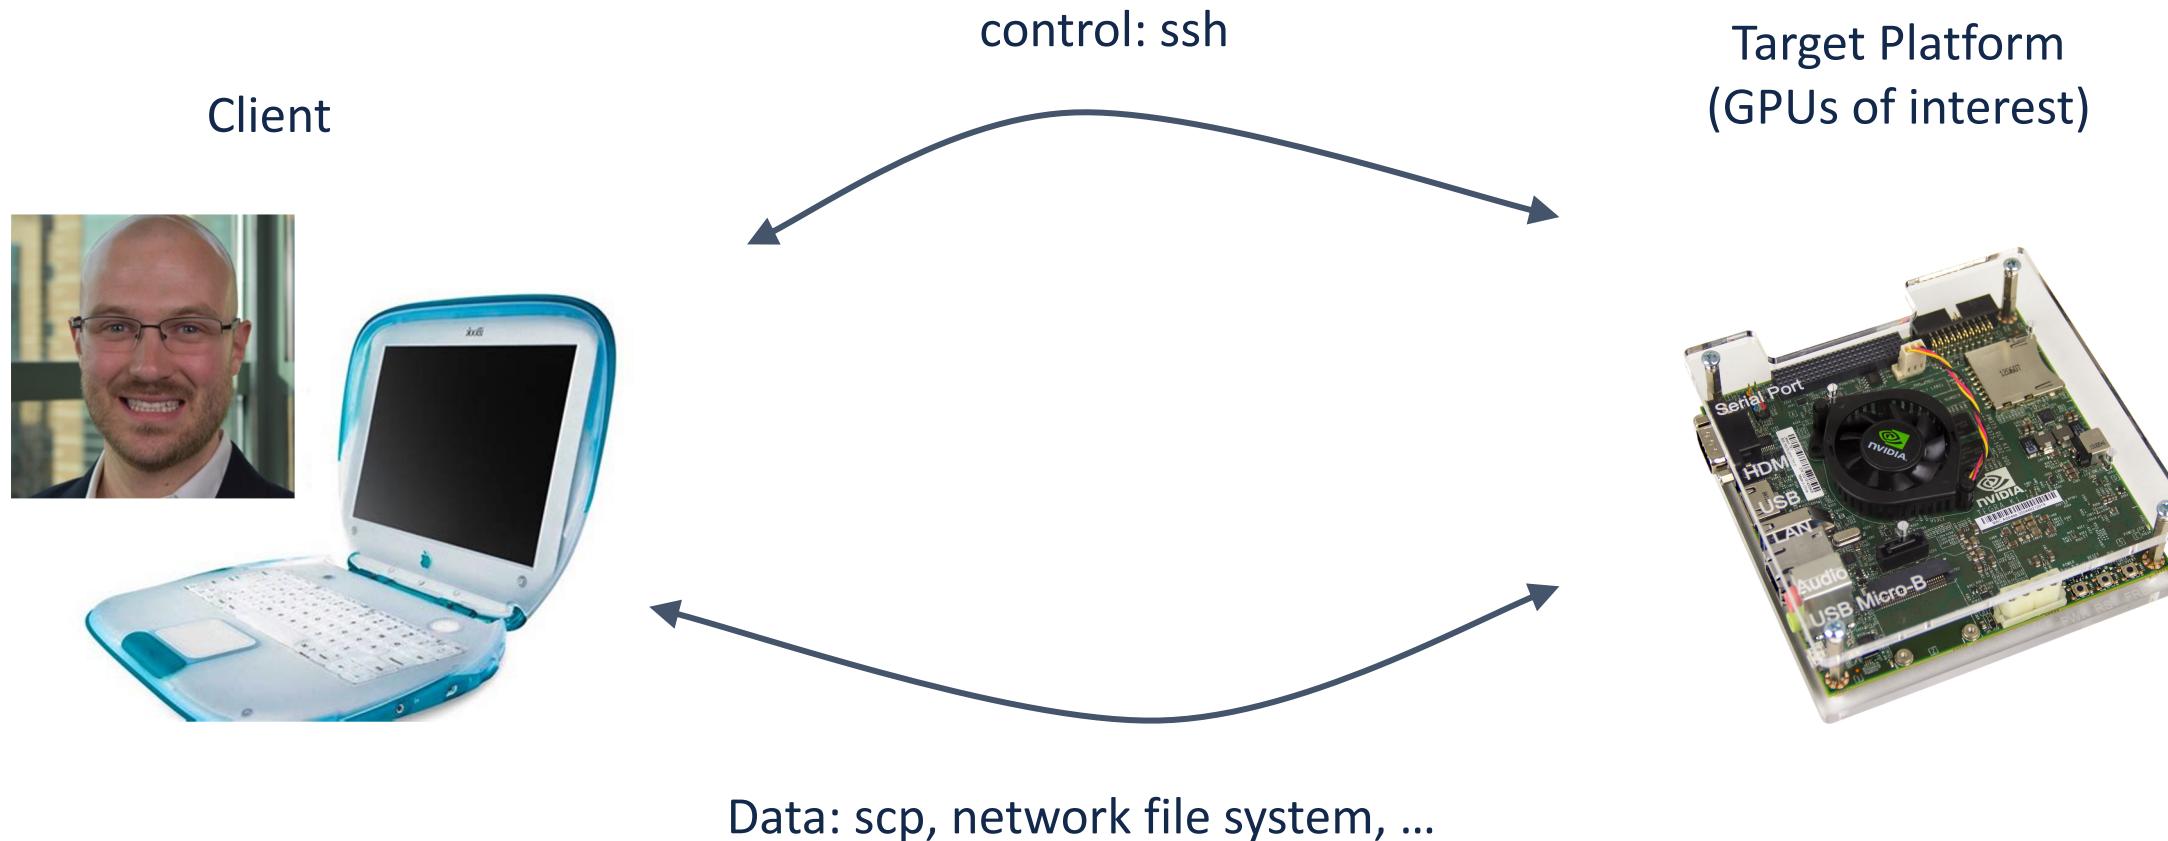
Resources

- Everything used in this lecture is at
github.com/cwpearson/nvidia-performance-tools
- Use it any way you want (with attribution).
 - Docker images for amd64 and ppc64le with CUDA and recent versions of Nsight
 - Matrix multiplication examples (in `sgemm/`)
 - `rai_build.yml` for `sgemm` (in `sgemm/`) if you have access to rai
 - Build and profile the examples on any system
- Profile data files and source code can be downloaded here:
<https://uofi.box.com/s/kh7plqoxqyy2yp54abokhqdttopga0oxh>

System- and Kernel-Level Profiling

- Nsight Compute: Kernel-Level Profiling
 - How fast does the GPU execute my kernel?
- Nsight Systems: System-level Profiling
 - How effectively is my system delivering work to the GPU?
 - What is my system doing while the GPU is working?
 - How fast is data moving to/from the GPU?
 - How much time does the CPU take to control the GPU?
 - When do asynchronous operations occur?

Common GPU Development Model



Record-and-analyze



Client

Target Platform
(GPUs of interest)



2

Transfer recorded files

3

Analyze with GUI
nsight-sys ...
nv-nsight-cu ...

1

Record with CLI
nsys profile ...
nv-nsight-cu-cli ...



Preparing for Profiling: Correctness

- Subtle errors that do not cause your kernel to terminate under normal conditions can cause errors with profiling
 - esp. writing outside of allocated memory
- Run your code with `cuda-memcheck` if profiling crashes or misbehaves
 - Automatically instruments for bad memory behavior
 - Causes something like 100x slowdown, so try small datasets first
 - Fix any errors that come up, then profile again

```
cuda-memcheck ./my-cuda-binary
```

Preparing for Profiling: Host Code Annotations

Nvidia Tools Extensions

#include <nvToolsExt.h> and link with -lnvToolsExt

Will show up as a named span in the Nsight System GUI timeline

Useful for marking parts of the code for later reference.

```
nvtxRangePush("sleeping");  
sleep(100);  
nvtxRangePop();
```

Preparing for Profiling: Compiling

- Compile device code *with optimizations*
 - non-optimized or debug code often has many more memory references
 - nvcc by default applies many optimizations to device code
 - remove any -G flag (this flag generates debug info for device code)
- Compile device code with line number annotations
 - add -lineinfo flag to all nvcc calls
 - puts some info in the binary about what source file locations generated what machine code

```
$ nvcc -G      →      $ nvcc -lineinfo  
main.cu          main.cu
```

Preparing for Profiling: Compiling

--generate-line-info / -lineinfo

Generate line-number information for device code.

Annotates the binary with information to correlate ptx back to CUDA source code

Compiled PTX

```
.loc 1 18 12 // file 1 line 18 col 12
cvta.to.global.u64 %rd1, %rd6;
mov.u32 %r27, %ctaid.x;
mov.u32 %r1, %ntid.x;
mov.u32 %r28, %tid.x;
mad.lo.s32 %r2, %r27, %r1, %r28;
```

CUDA Source Code

```
18: int gidx = blockDim.x *
19:   blockIdx.x + threadIdx.x;
```

Preparing for Profiling: Compiling

Don't use any of these for Nsight profiling!

--profile / -pg

Instrument generated code/executable for use by gprof (Linux only).

--debug / -g

Generate debug information for host code.

--device-debug / -G

Generate debug information for device code. Turns off all optimizations.

Don't use for profiling; use -lineinfo instead.

Preparing for Profiling: System

Nsight System uses various system hooks to accomplish profiling.

Some errors would reduce the amount or accuracy of gathered info, some will make system profiling impossible. Consult the documentation for how to correct.

An example of a GOOD output: (check with `nsys status -e`)

```
$ nsys status -e
Sampling Environment Check
Linux Kernel Paranoid Level = 2: OK
Linux Distribution = Ubuntu
Linux Kernel Version = 4.16.15-41615: OK
Linux perf_event_open syscall available: OK
Sampling trigger event available: OK
Intel(c) Last Branch Record support: Available
Sampling Environment: OK
```

Caveats

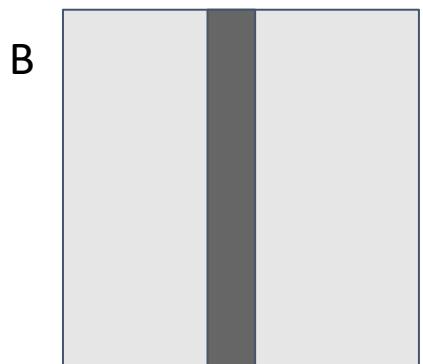
Profiling affects the performance of your kernel!

It will help you improve the speed, but do not report the time *during* profiling as the performance of your code. Always run and time without profiling.

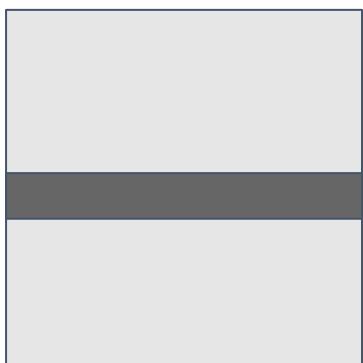
Matrix Multiplication Review

Dense Matrix Multiplication

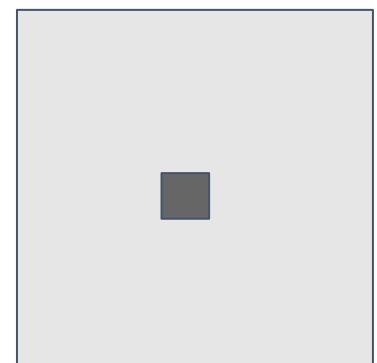
$$C = A \times B$$



A

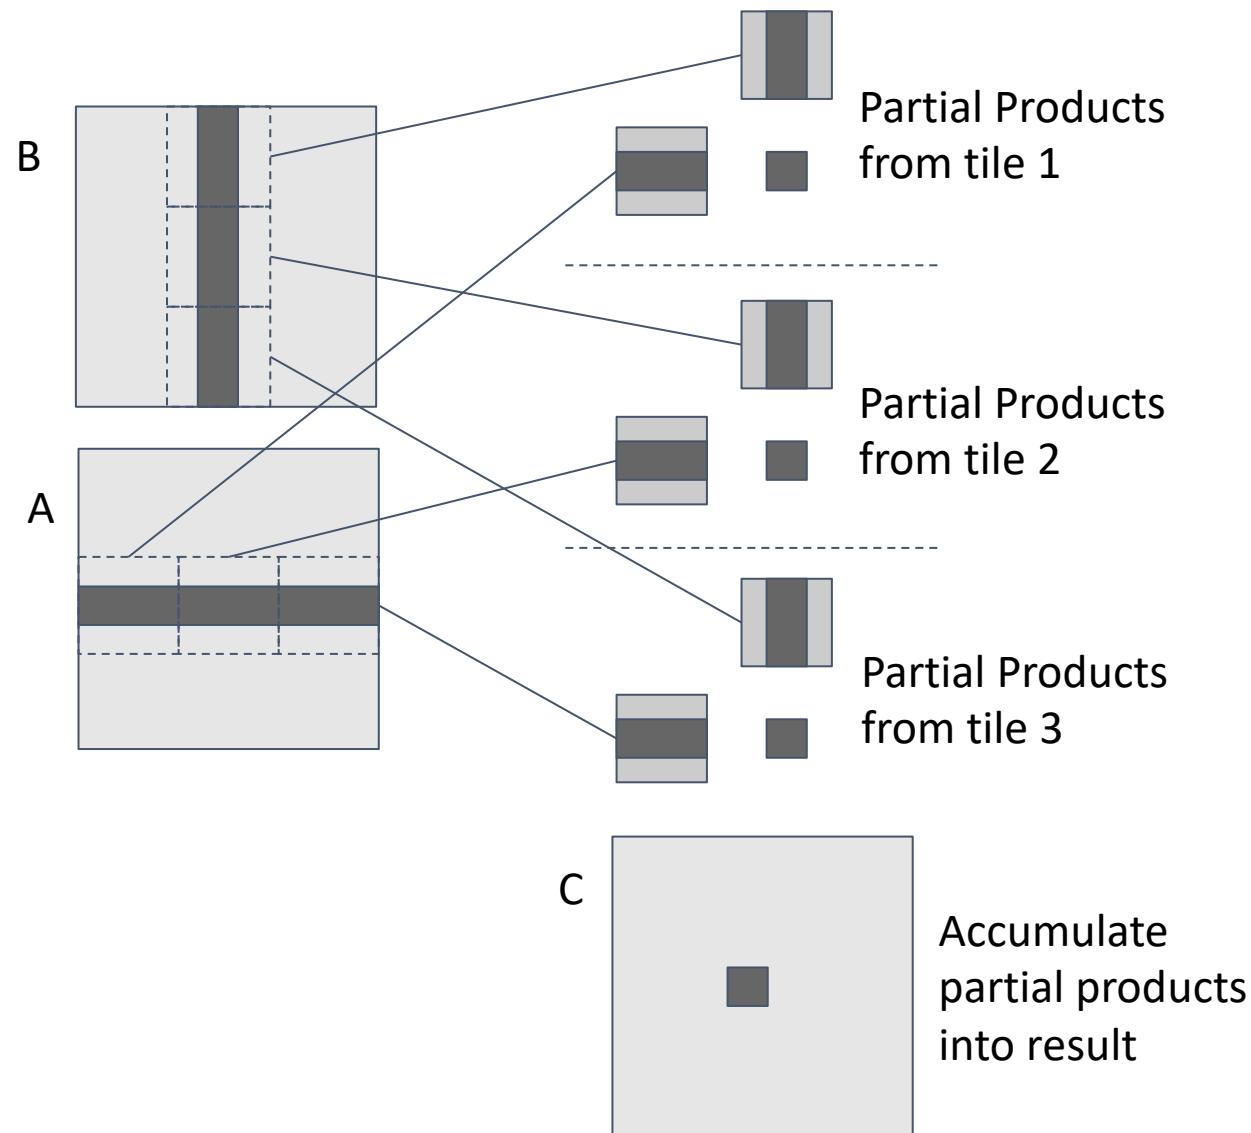


C



- Each thread produces a single product value $C_{i,j}$ by $\text{dot}(A_i, B_j)$
- A and C are column-major, B is row-major
 - access to B is coalesced
- Each entry of the A/B matrices loaded from global memory multiple times

Shared-Memory Tiling



- Each thread produces a single product value $C_{i,j}$ by $\text{dot}(A_i, B_j)$
- Each thread block collaboratively loads tiles of A and B to accumulate partial products
 - Much reuse comes from fast shared memory instead of slow global memory

SGEMM Comparison

	A Reuse	B Reuse	Product Data per Block	Shared Memory/block
Basic	1	1	1024	0
Tiled	32 (TILE_SIZE)	32 (TILE_SIZE)	1024 (TILE_SIZE ²)	32*32*2*4B = 8KB

Measuring Time with CUDA Events

Timing Operations

```
#include <chrono>
typedef std::chrono::high_resolution_clock Clock;
typedef std::chrono::duration<float> Duration;
cudaDeviceSynchronize();
auto wallStart = Clock::now();
kernel<<<...>>>(...);
cudaDeviceSynchronize();
Duration wallElapsed = Clock::now() - wallStart;
std::cout << wallElapsed.count() << " s";
```

Ensure the GPU is not doing other work.

Get wall time **before** and **after** the things you want to measure.

Be sure to synchronize device so asynchronous operations have finished.

- cudaDeviceSynchronize() may be relatively costly for fast operations
- Must synchronize with CPU
- CPU may be needed for other things during GPU execution

Terminology

- Stream (`cudaStream_t`)
 - A queue of sequential CUDA events. Each is executed after the prior one finishes
 - A program can use any number of CUDA streams
 - Associated with a device
- Default Stream (`cudaStream_t stream = 0`)
 - A special stream that is used when no stream is provided
- Event (`cudaEvent_t`)
 - Records the state of a stream
- See CUDA programming guide for stream synchronization edge cases
- *Generally*, to overlap operations:
 - different streams
 - do not use pageable memory
 - use `*async` CUDA runtime functions

Timing Async Operations with CUDA Events

```
cudaEvent_t start, stop;  
cudaStream_t stream;  
cudaStreamCreate(&stream);  
cudaEventCreate(&start);  
cudaEventCreate(&stop);  
cudaEventRecord(start, stream);  
cudaMemcpyAsync(..., stream);  
kernel<<<dimGrid, dimBlock, 0, stream>>>();  
cudaEventRecord(stop, stream);  
cudaMemcpyAsync(..., stream);  
cudaEventSynchronize(stop);  
float millis;  
cudaEventElapsedTime(&millis, start, stop);
```

Place and events in the stream **before** and **after** the **things you want to measure**. Executed when the stream reaches the event, not when cudaEventRecord is called.

Wait for the final event to be reached. Could use cudaStreamSynchronize or cudaDeviceSynchronize too.

Get the time between the start and stop event.

Warmup & Multiple Measurements

```
* Running bash -c "./1-1-pinned-basic | tee 1-1-pinned-basic.txt"
generate data
0: 3.75206
1: 3.73158
2: 3.73213
3: 3.73149
4: 3.73379
5: 3.73146 *
6: 3.73043 *
7: 3.72659 *
8: 3.73094 *
9: 3.72835 *
kernel 1787.02GFL0PS (6664784846 flop, 0.00372956s)
```

warmup runs

these contribute to reported time

average kernel performance

Kernel Profiling with Nsight Compute

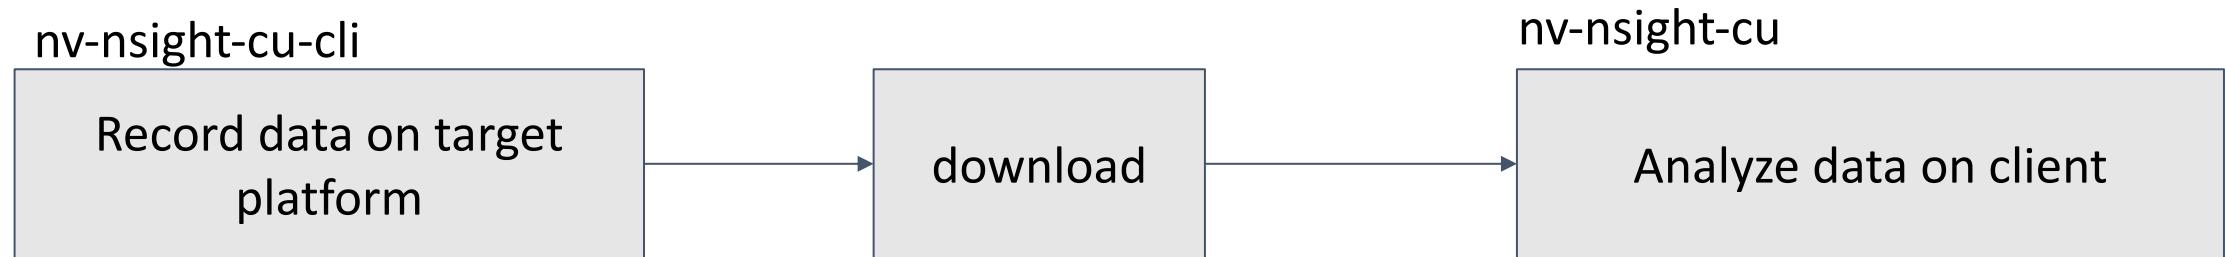
Problem Setup

- Dense matrix $C = A * B$
- C and A are column-major, B is row-major
- All matrices 1500×1500 floats
 - 9 MB each, 27 MB total (larger than L2 cache)

Kernel	Performance	Speedup
Basic	1787 GFLOPS	-
Tiled	2585 GFLOPS	1.45

Nvidia Nsight Compute

- Record and analyze detailed kernel performance metrics
- Two interfaces:
 - GUI (nv-nsight-cu)
 - CLI (nv-nsight-cu-cli)
- Directly consuming 1000 metrics is challenging, we use the GUI to help



Kernel Profiling

- Device has many performance counters to record detailed information
 - Made available as “metrics”.
 - Titan V on rai supports ~1100 metrics, some shown below
 - `$ nv-nsight-cu-cli --devices 0 --query-metrics`

<code>lts__t_sectors_srcunit_l1_op_atom_dot_cas</code>	# of LTS sectors from unit L1 for atomic CAS
<code>l1tex__data_pipe_lsu_wavefronts_mem_shared_cmd_write</code>	# of shared write wavefronts processed by Data-Stage
<code>lts__t_sectors_srcunit_l1_aperture_sysmem_op_read</code>	# of LTS sectors from unit L1 accessing system memory
<code>(sysmem) for reads</code>	
<code>lts__t_requests_op_red_lookup_hit</code>	# of LTS requests for reductions that hit
<code>lts__t_sectors_equiv_l1tagmiss_pipe_tex_mem_texture_op_id</code>	# of sectors requested for TLD instructions
<code>l1tex__t_bytes_pipe_tex_lookup_miss</code>	# of bytes requested that missed for TEX pipe
<code>l1tex__texin_requests_mem_texture</code>	# of texture requests (quads) sent to
<code>TEXINI1tex__t_bytes_pipe_lsu_mem_local_op_id_lookup_miss</code>	# of bytes requested that missed for local
<code>loadsl1tex__t_bytes_pipe_tex_mem_surface_op_red_lookup_miss</code>	# of bytes requested that missed for surface reductions
<code>...</code>	

Record Kernel Trace

```
$ nv-nsight-cu-cli \
  --kernel-id ::mygemm:6 \
  --section ".*" \
  -o 1-1-pinned-basic \
  1-1-pinned-basic
```

Profile the 6th time the “mygemm” kernel runs
Record metrics for all report sections
Create “1-1-pinned-basic.nsight-cuprof-report”
Name of the CUDA executable to profile

Nsight Compute Sections

A group of related measurements

The default list can be generated by

```
$ nv-nsight-cu-cli --list-sections
```

Without the --sections options, this is what would be recorded

We provide a regex that matches all sections

Identifier	Display Name	Enabled	Filename
ComputeWorkloadAnalysis	Compute Workload Analysis	no	.../.../.../sections/ComputeWorkloadAnalysis.section
InstructionStats	Instruction Statistics	no	...64/.../.../sections/InstructionStatistics.section
LaunchStats	Launch Statistics	yes	...1_3-x64/.../.../sections/LaunchStatistics.section
MemoryWorkloadAnalysis	Memory Workload Analysis	no	...4/.../.../sections/MemoryWorkloadAnalysis.section
MemoryWorkloadAnalysis_Chart	Memory Workload Analysis Chart	no/sections/MemoryWorkloadAnalysis_Chart.section
MemoryWorkloadAnalysis_Tables	Memory Workload Analysis Tables	no	.../sections/MemoryWorkloadAnalysis_Tables.section
Occupancy	Occupancy	yes	...ibc_2_11_3-x64/.../.../sections/Occupancy.section
SchedulerStats	Scheduler Statistics	no	...-x64/.../.../sections/SchedulerStatistics.section
SourceCounters	Source Counters	no	..._11_3-x64/.../.../sections/SourceCounters.section
SpeedOfLight	GPU Speed Of Light	yes	..._2_11_3-x64/.../.../sections/SpeedOfLight.section
WarpStateStats	Warp State Statistics	no	...-x64/.../.../sections/WarpStateStatistics.section

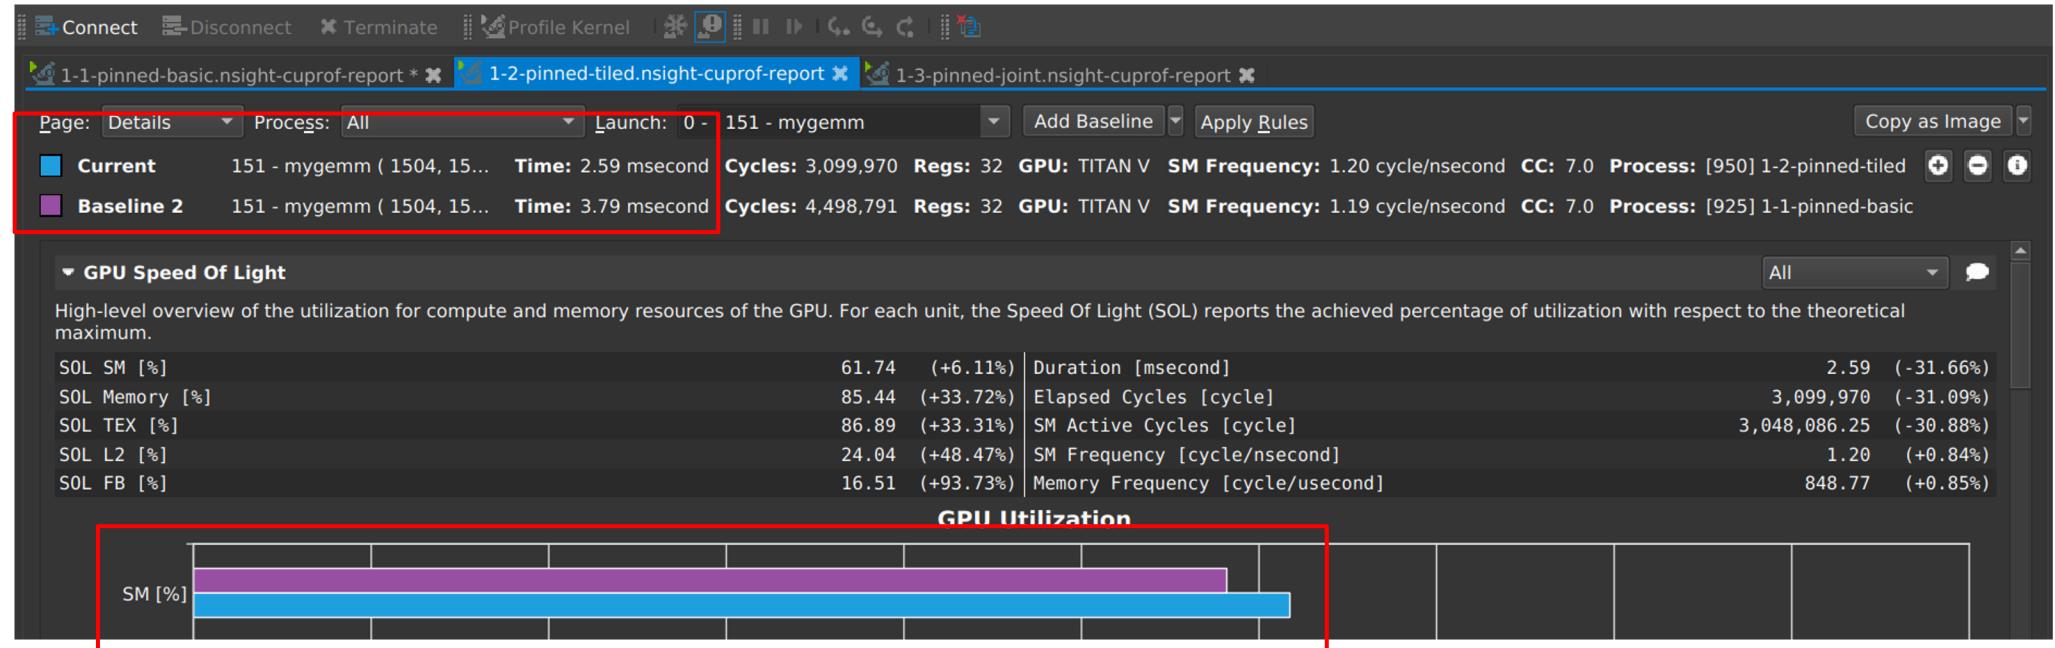
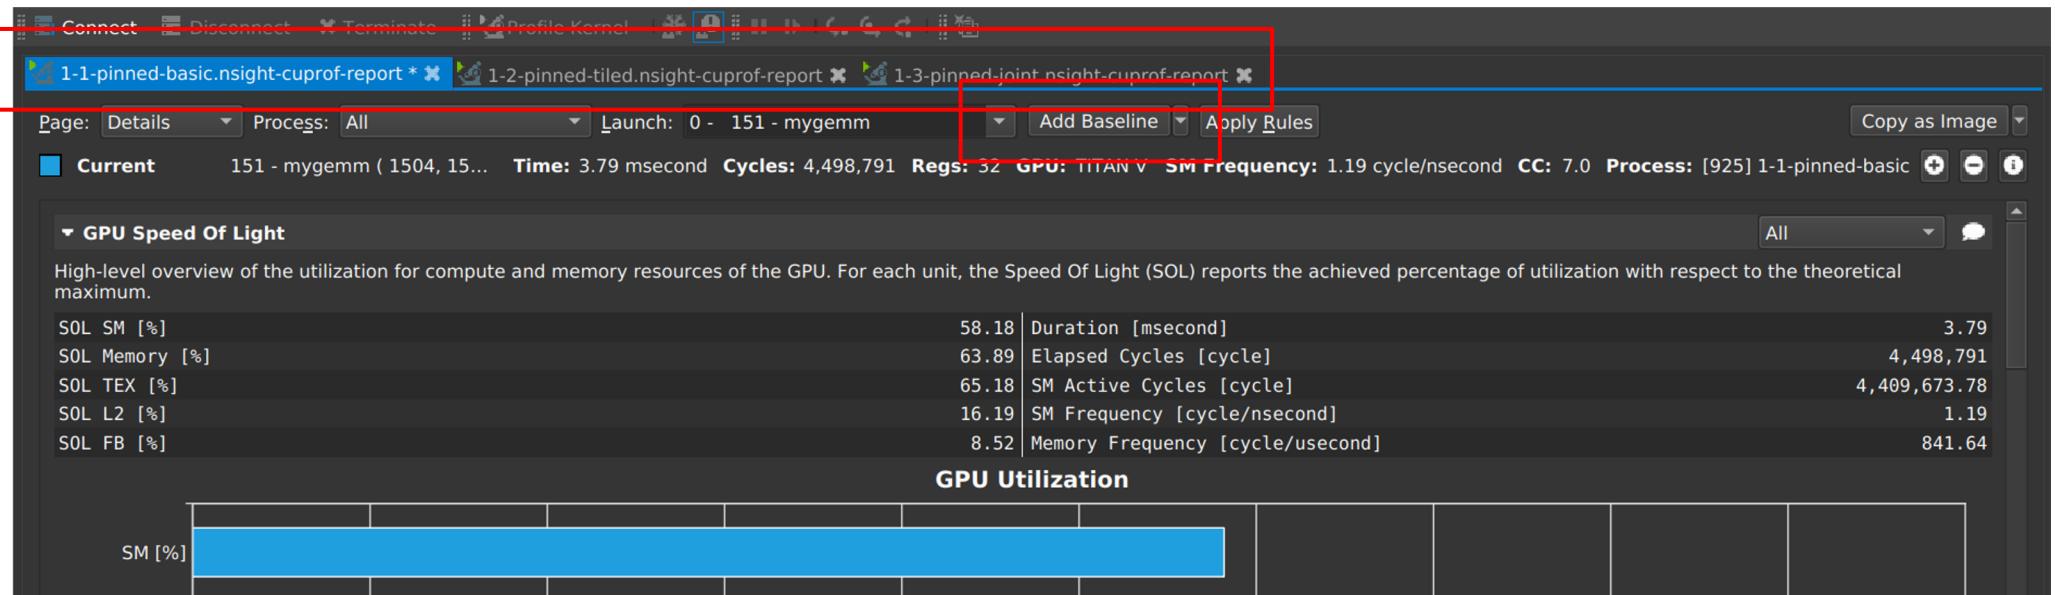
Open in Nsight Compute

Start Nsight Compute

File > Open File ... > 1-1-pinned-basic.nsight-cuprof-report

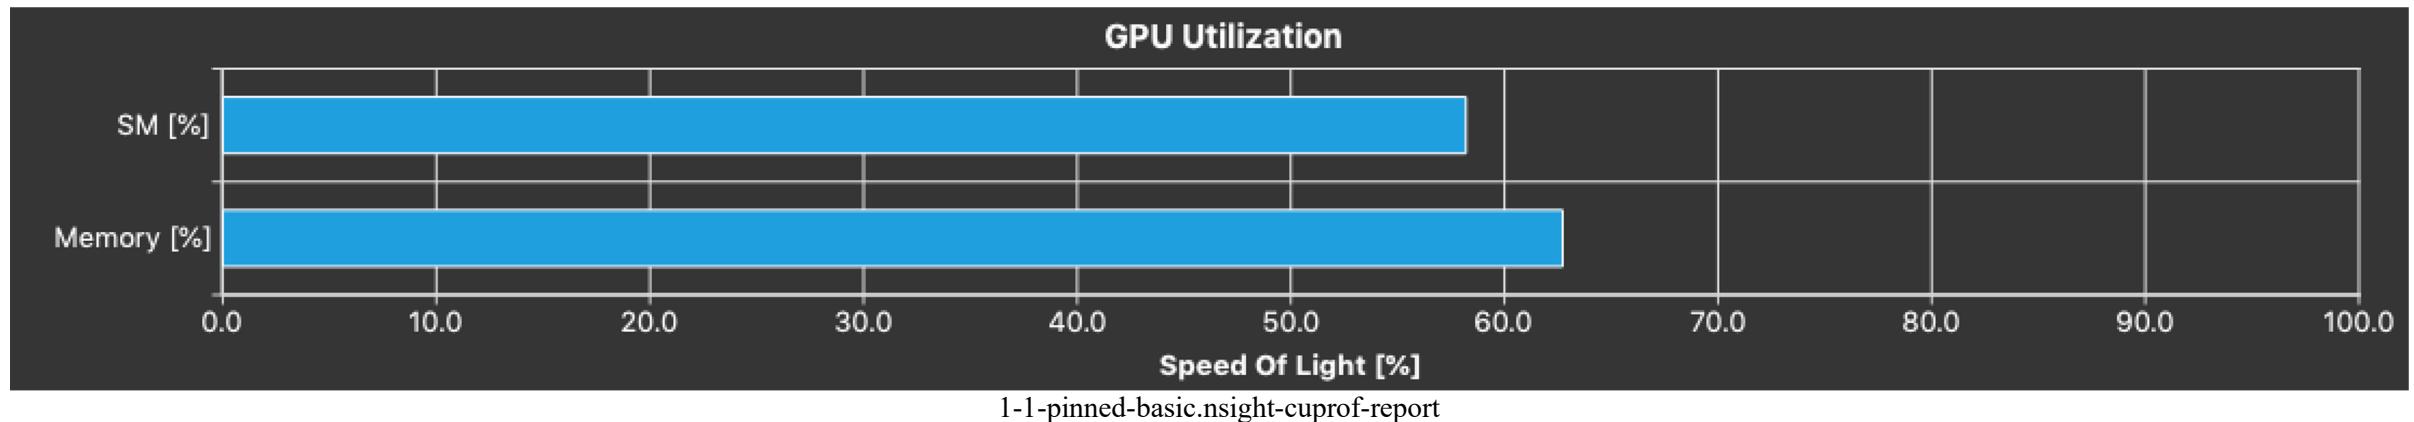
- Can open multiple files, will be open in multiple tabs
 - Can also use different runs as “baselines” for comparison in the same tab
 - Click “Add Baseline”

Tabs and baseline button



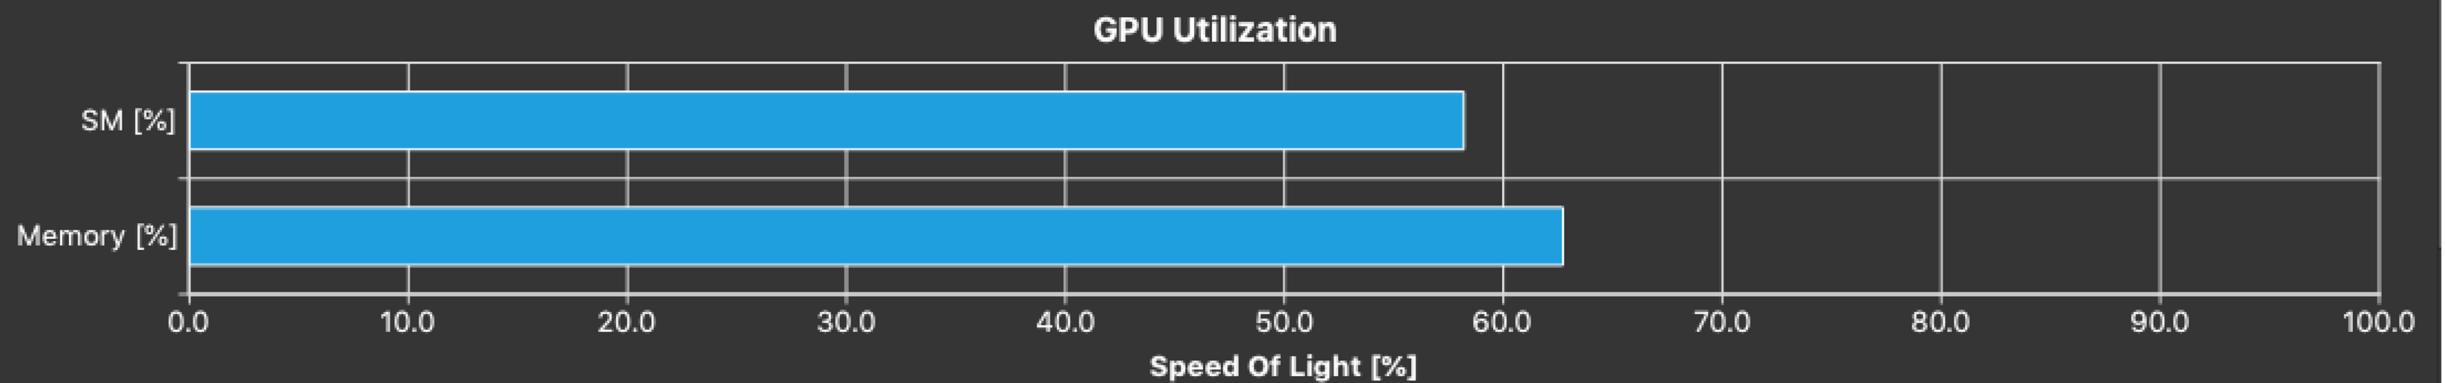
Next tab now has comparison

GPU “Speed of Light” (SoL): High-Level Bottleneck



Achieved % of utilization with respect to maximum

	SM < 60	SM > 60
Memory < 60	Latency	Compute
Memory > 60	Bandwidth	Compute & Bandwidth



SOL SM Breakdown

SOL SM: Inst Executed Pipe Lsu [%]	58.20
SOL SM: Issue Active [%]	52.98
SOL SM: Inst Executed [%]	52.98
SOL SM: Pipe Fma Cycles Active [%]	39.23
SOL SM: Pipe Alu Cycles Active [%]	36.56

Mouse over to
see metric
name

SOL Memory Breakdown

SOL L1: Data Pipe Lsu Wavefronts [%]	62.72
SOL L1: Lsu Writeback Active [%]	62.45
SOL L1: Lsuin Requests [%]	58.20
SOL L2: T Sectors [%]	16.60
SOL L2: Lts2xbar Cycles Active [%]	15.50

**SM is mostly busy issuing memory operations.
58.2% instructions executed by Load-Store Unit.
only 35-40% of cycles ALU / FMA is busy**

Section: GPU Speed of Light

- Achieved percentage of utilization w.r.t theoretical maximum

	Basic	Tiled
(GFLOPS)	1787	2585
SoL SM	58.18	61.74
SoL Memory	63.89	85.44

Workload Memory Analysis: Memory Chart

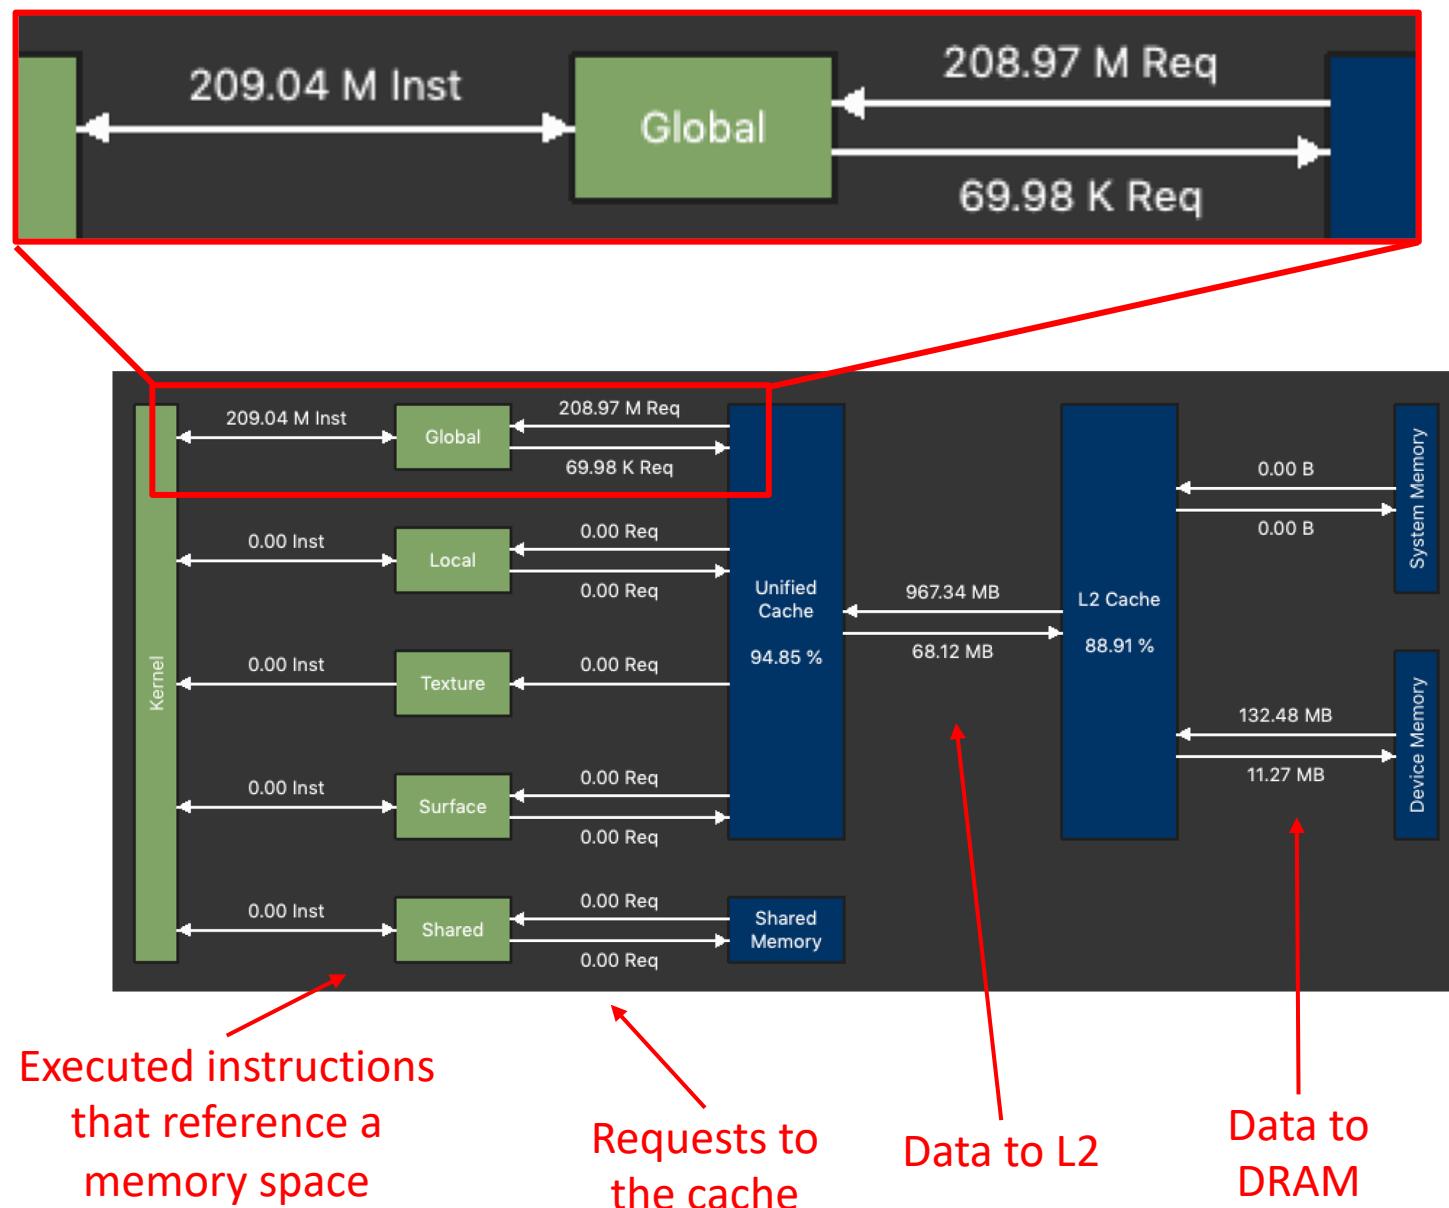
Global Memory: shared by all threads

Local Memory: private per-thread

Shared Memory: shared by threads in a block

Texture/Surface: Cached for 2D spatial locality

Constant (?): Cached in the constant cache



Memory Workload Analysis: Charts

Shared Memory									
	Instructions	Requests	% Peak	Bank Conflicts					
Shared Load	26,999,808	54,031,933	52.77	30,783					
Shared Store	1,687,488	1,941,325	1.90	253,837					
Shared Atomic	0	-	-	-					
Total	28,687,296	55,973,258	54.67	284,620					
First-Level (Unified) Cache									
	Instructions	SM->TEX Requests	% Peak	Hit Rate	TEX->L2 Requests	% Peak	L2->TEX Returns	% Peak	TEX->SM Returns
Global Load Uncached	-	-	-	-	-	-	31,751,521	31.01	64,551,202
Local Load Cached	0	0	0	0	-	-	-	-	-
Local Load Uncached	-	-	-	-	-	-	-	-	-
Surface Load	0	0	0	0	-	-	0	0	0
Texture Load	0	0	0	0	-	-	0	0	0
Global Store	70,453	70,453	0.07	25.03	340,619	0.33	-	-	-
Local Store	0	0	0	0	-	-	-	-	-
Surface Store	0	0	0	0	-	-	-	-	-
Global Reduction	0	0	0	0	-	-	-	-	-
Surface Reduction	0	0	0	0	-	-	-	-	-
Global Atomic	0	0	0	0	-	-	-	-	-
Global Atomic Cas	0	0	0	0	-	-	-	-	-
Surface Atomic	0	0	0	0	-	-	-	-	-
Surface Atomic Cas	0	0	0	0	-	-	-	-	-
Loads	8,281,306	8,281,306	8.09	25.38	-	-	31,751,521	31.01	64,551,202
Stores	70,453	70,453	0.07	25.03	340,619	0.33	-	-	-
Total	8,351,759	8,351,759	8.16	25.38	340,619	0.33	31,751,521	31.01	64,551,202
Second-Level (L2) Cache									
	TEX->L2 Requests	% Peak	L2->TEX Returns	% Peak	Total Bytes	Total Throughput			
Global Load Cached	-	-	-	-	1,016,048,672	948,712,830,166.13	-	-	
Global Load Uncached	-	-	31,751,521	31.01	-	-	-	-	
Local Load Cached	-	-	-	-	-	-	-	-	
Local Load Uncached	-	-	-	-	-	-	-	-	
Surface Load	-	-	0	0	0	0	0	0	
Texture Load	-	-	0	0	0	0	0	0	
Global Store	340,619	0.33	-	-	10,899,808	10,177,453,089.52	-	-	
Local Store	-	-	-	-	-	-	-	-	
Surface Store	0	0	-	-	0	0	0	0	
Global Reduction	0	0	-	-	0	0	0	0	
Surface Reduction	0	0	-	-	0	0	0	0	
Global Atomic	0	0	-	-	0	0	0	0	
Global Atomic Cas	0	0	-	-	0	0	0	0	
Surface Atomic	0	0	-	-	0	0	0	0	
Surface Atomic Cas	0	0	-	-	0	0	0	0	
Loads	-	-	31,751,521	31.01	1,016,048,672	948,712,830,166.13	-	-	
Stores	340,619	0.33	-	-	10,899,808	10,177,453,089.52	-	-	
Total	340,619	0.33	31,751,521	31.01	1,026,948,480	958,890,283,255.65	-	-	
Device Memory (FB)									
	L2<->FB Sectors	% Peak	Bytes	Throughput					
Load	857,924	3.92	27,453,568	25,634,158,001.67	-	-	-	-	
Store	313,103	1.43	10,019,296	9,355,294,609.78	-	-	-	-	
Total	1,171,027	5.35	37,472,864	34,989,452,611.45	-	-	-	-	

- Detailed information summarized in the Memory Chart
- Uses TEX to mean the first-level cache.

Memory Workload Analysis

	Basic	Tiled
GFLOPS	1787	2585
Speed of Light: Memory	63.89	85.44
Global Load Cached (% peak)	59.24	2.70
Global Load Cached (SM->TEX REQ)	209M	6.6M
Shared Load (REQ)	0	160M
L1 Hit Rate	94.84	74.75
Global Load (B)	197M	266M
Global Store (B)	11.7M	11.9M

Replaced global loads with shared loads

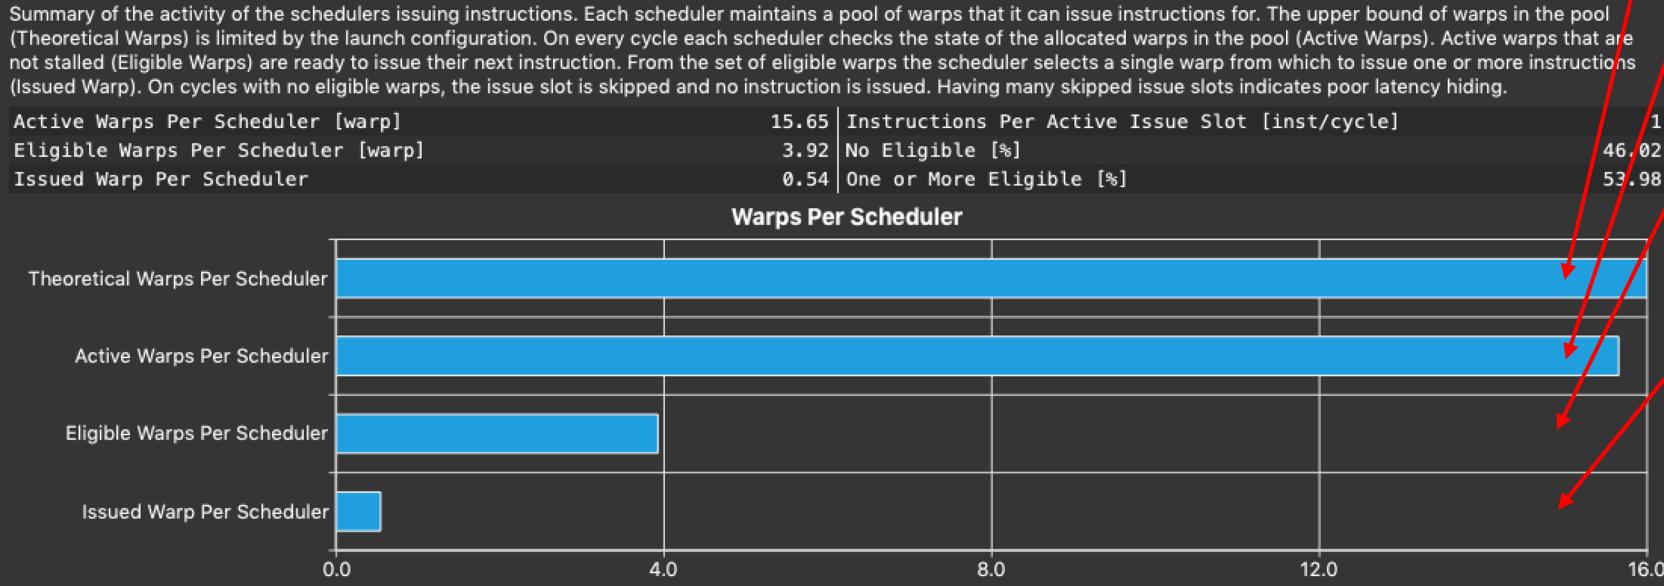
Cache reduces global memory traffic better than shared memory?

Scheduler Statistics

	Maximum Fermi	Maximum Kepler	Scope	Limiter
Device Limit	48 Warps	64 Warps		
Theoretical Occupancy	≤ 48 Warps	≤ 64 Warps	Launch	Launch Configuration
Active Warps			Warp Launch & Complete	Scheduler Load Balancing
			Cycle	
			Cycle	Issue Stalls
			Cycle	Eligible Warps & # Warp Schedulers
Stalled Warps				
Eligible Warps				
Selected Warps	≤ 2 Warps	≤ 4 Warps		

Scheduler Statistics

Scheduler Statistics !



Pool of warps that the scheduler can pick from. Limited by device.

Number of warps actually given to SM: not enough work, or work imbalance

Number of warps ready to execute: waiting for barrier, watching for instruction fetch, waiting for data...

Number of issued warps: usually maximum of 1 or 2 depending on hardware.

Just because average value is good, doesn't mean warp scheduling chances are missed

Scheduler Statistics Comparison

	Basic	Tiled
GFLOPS	1787	2585
Theoretical Warps / Scheduler	16	16
Active Warps / Scheduler	15.65	15.70
Eligible Warps / Scheduler	3.92	2.58
Issued Warps / Scheduler	0.54	0.33
No Eligible (%)	46	67
One or More Eligible (%)	54	33

Average 2-3 cycles
between a warp
being ready

Warp State Statistics

Warp State Statistics

Analysis of the states in which all warps spent cycles during the kernel execution. The warp states describe a warp's readiness or inability to issue its next instruction. The warp cycles per instruction define the latency between two consecutive instructions. The higher the value, the more warp parallelism is required to hide this latency. For each warp state, the chart shows the average number of cycles spent in that state per issued instruction. Stalls are not always impacting the overall performance nor are they completely avoidable. Only focus on stall reasons if the schedulers fail to issue every cycle.

Warp Cycles Per Issued Instruction	29.00	Avg. Active Threads Per Warp	31.89
Warp Cycles Per Issue Active	29.00	Avg. Not Predicated Off Threads Per Warp	31.86
Warp Cycles Per Executed Instruction [cycle]	29.00	-	-

Avg. Active Threads Per Warp

Avg. Not Predicated Off Threads Per Warp

latency between two consecutive instructions

More latency: more warp parallelism needed
to hide

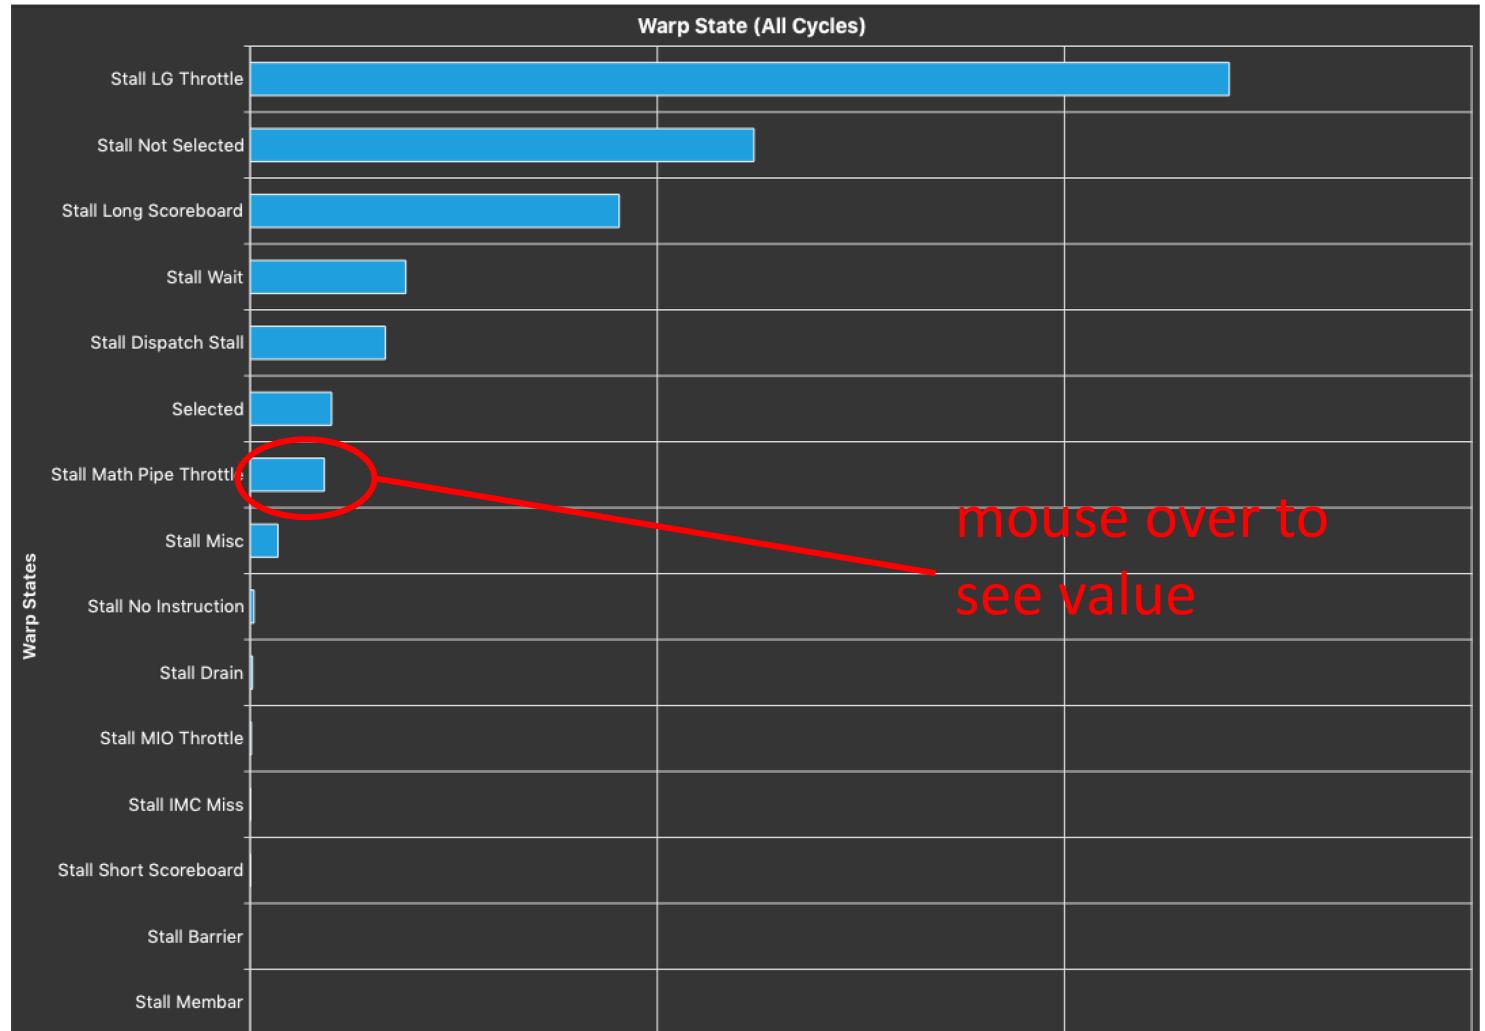
Block size not multiple of warp size
or
threads do not exit kernel at same time

Average control
divergence

Warp State Statistics

What portion of the cycles per instruction the warps were stalled in each state.

Stalls cannot always be avoided and only really matter if instructions can't be issued every cycle



Warp State Statistics

	Description	Basic	Tiled
GFLOPS		1787	2585
Warp Cycles per Issued Instruction	latency between two consecutive instructions	29.00	47.34
Stall Long Scoreboard	Waiting for local, global, texture, or surface load	4.53	6.52
Stall Barrier	Waiting at barrier	0	4.73
Stall MIO Throttle	Waiting for MIO queue, caused by loads (incl. shared), and special math	0.01	22.61
Stall LG Throttle	Waiting for load-store-unit queue for local and global memory instructions	11.83	1.79
Stall Not selected	Eligible but not selected because another eligible warp was	6.89	6.82

Replaced global loads with shared

Section: Occupancy

Occupancy

Occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however, low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.

Theoretical Occupancy [%]	100	Block Limit Registers [block]	2
Theoretical Active Warps per SM [warp/cycle]	64	Block Limit Shared Mem [block]	32
Achieved Occupancy [%]	97.77	Block Limit Warps [block]	2
Achieved Active Warps Per SM [warp]	62.57	Block Limit SM [block]	32

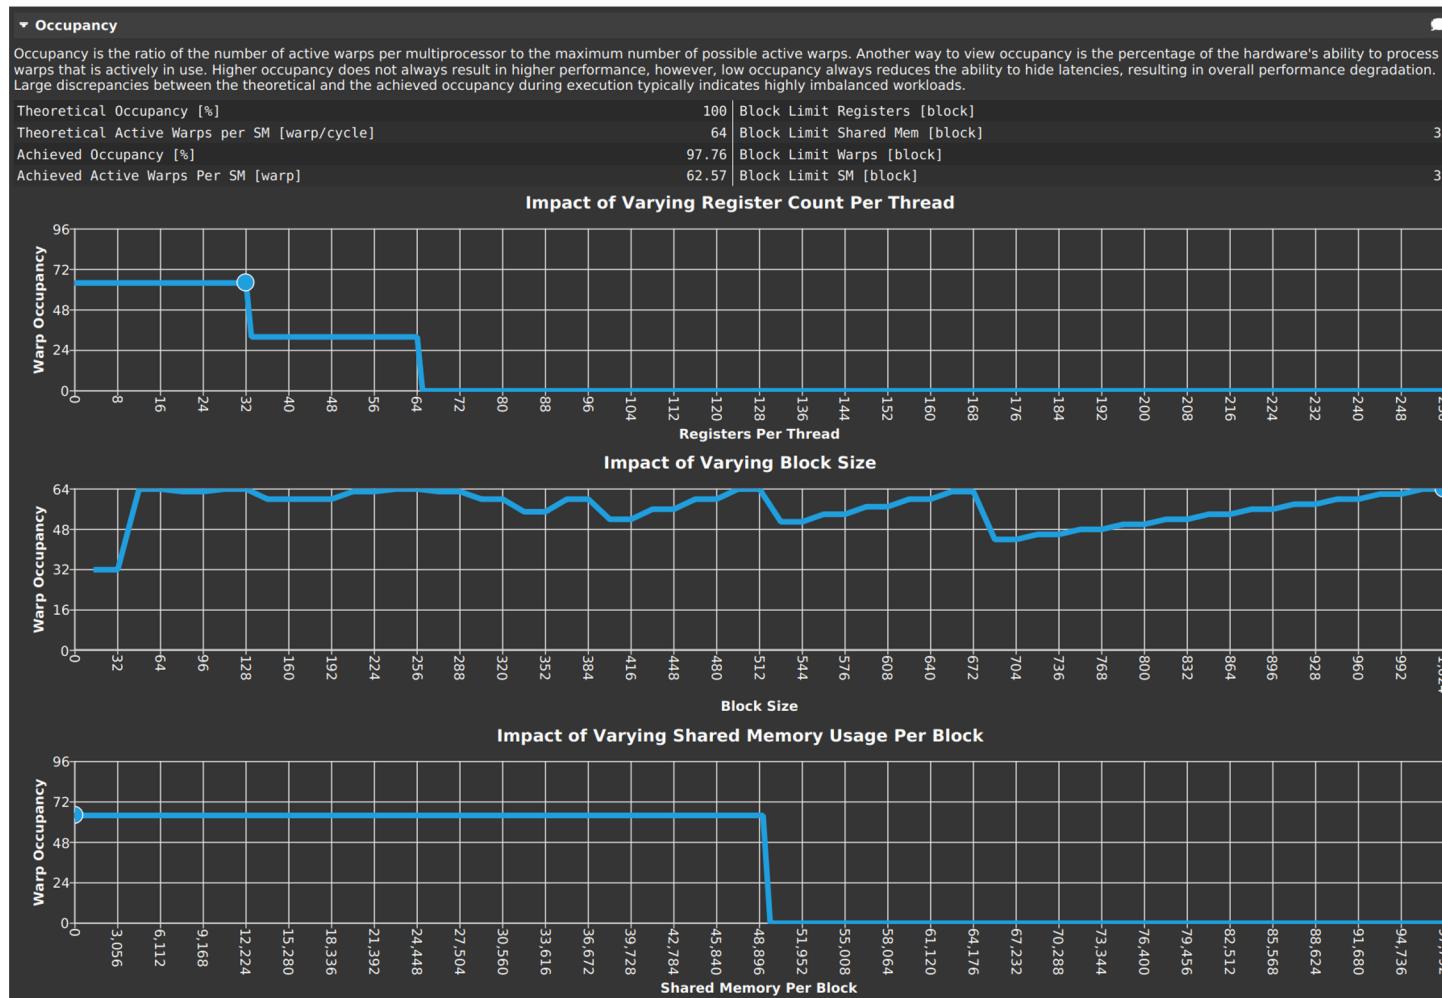
Achieved occupancy: true amount of average active warps.

Lower if workload is imbalanced across blocks, too few blocks, or last wave is empty

Each SM can handle fixed number of warps, blocks, registers, and shared memory.

How each requirement restricts the theoretical occupancy

Section: Occupancy



Theoretical occupancy limited by device hardware and launch configuration

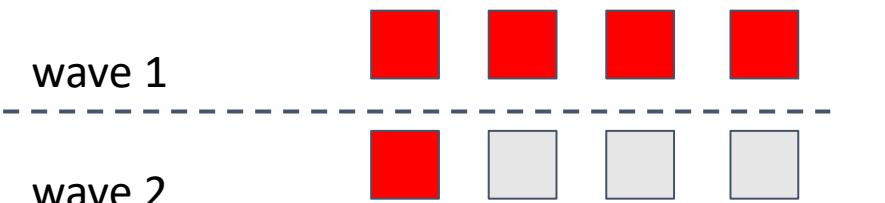
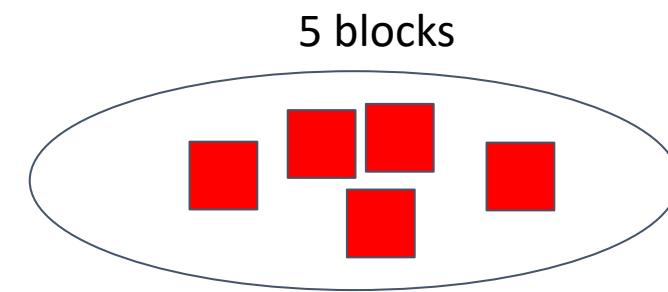
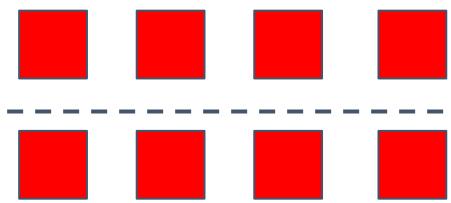
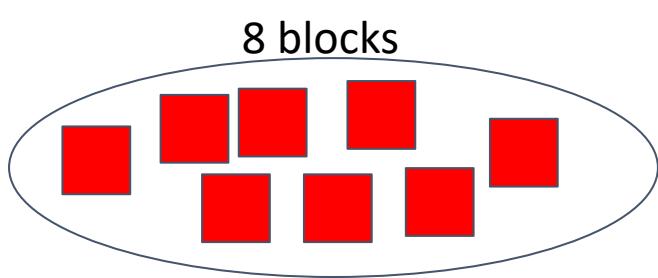
Achieved occupancy: true number of active warps as average

Lower if workload within or across blocks is imbalanced, if there are too few blocks, or the last wave is not large enough to fill GPU

Charts show how resources affect theoretical occupancy

Waves

Assume: 1 block per SM, GPU with 4 SMs



Launch Statistics and Occupancy

	Basic	Tiled
GFLOPS	1787	2585
Theoretical Occupancy	100	100
Th. Active Warps per SM	64	64
Achieved Occupancy	97.76	98.34
Waves per SM	13.81	13.81
Registers Per Thread	32	32

May not include registers for program counter!
Consult Nvidia's architecture whitepapers.
Titan V uses 2 additional registers for PC.

Instruction Hotspots

Show various metrics correlated with source code lines and PTX instructions

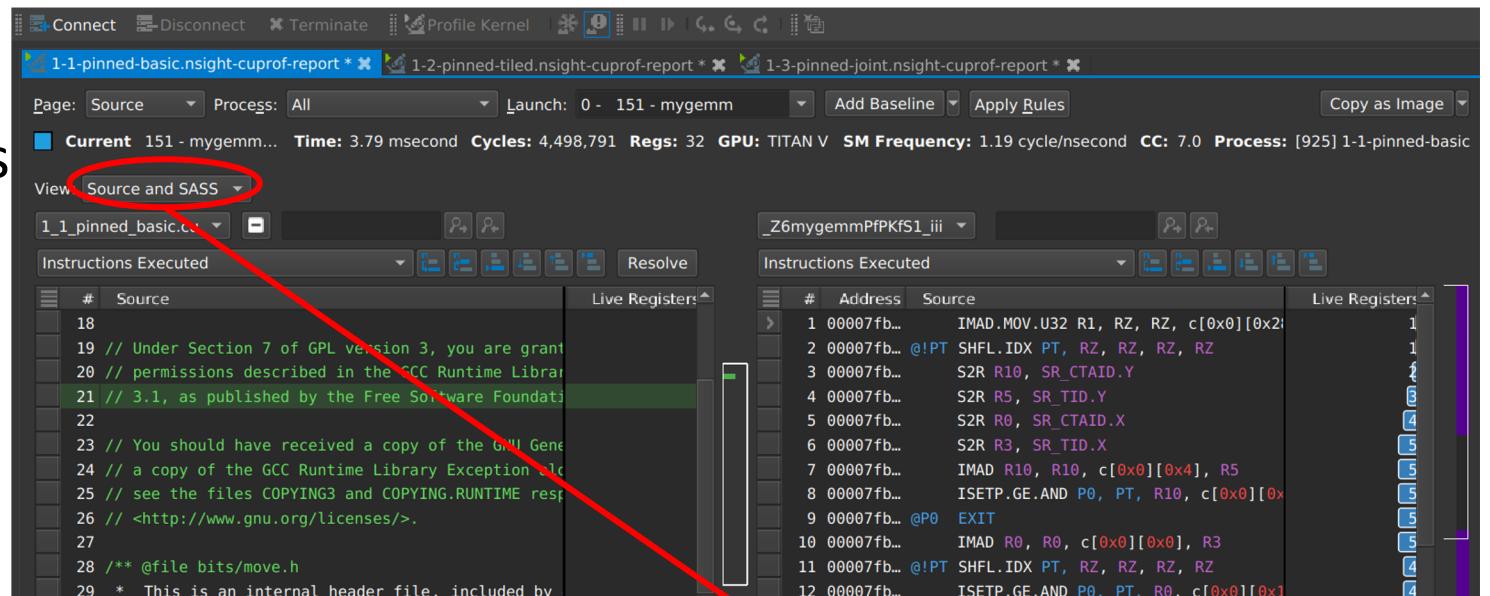
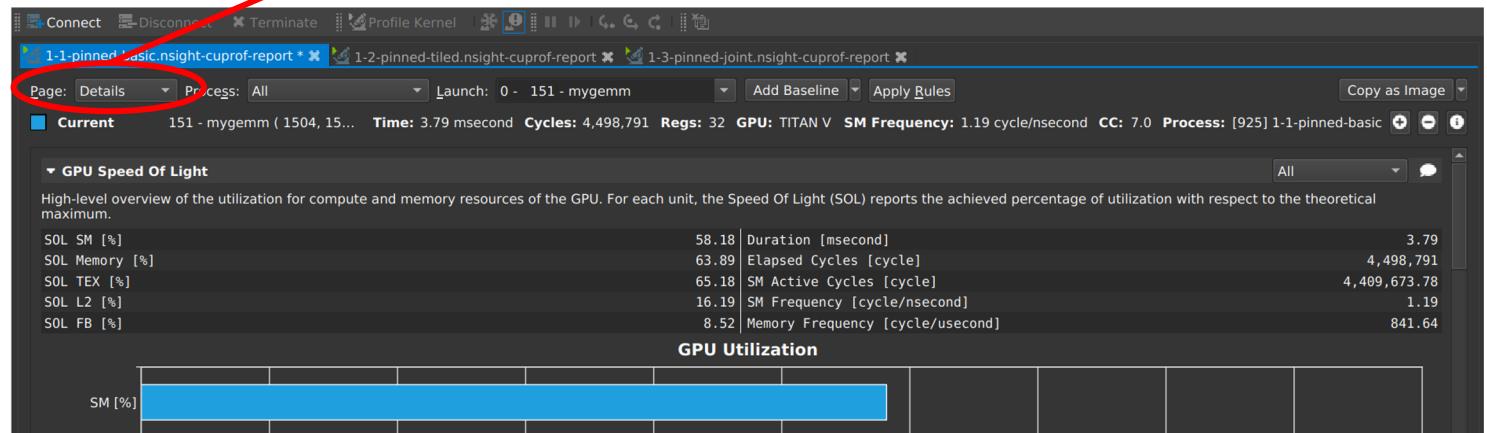
Some source code lines create many many PTX instructions: sometimes, split up a source line into many lines to get more details

```
dst[i] = src[i] + reg;
```

vs

```
temp v = src[i];
v += reg;
dst[i] = v;
```

Switch page to “Source”



“Source and PTX” (usually) or “Source and SASS”

PTX: higher-level assembly, same between GPU models

SASS: specific code for a GPU model

Instruction Hotspots

- If profiling on a different system, source file may not automatically load since paths may not match.

The screenshot shows the Nsight Compute interface with three tabs open: "1-1-pinned-basic.nsight-cuprof-report", "1-2-pinned-tiled.nsight-cuprof-report", and "1-3-pinned-joint.nsight-cuprof-report". The first tab is selected. The main window displays performance metrics: Current process 151 - mygemm..., Time: 3.79 msecond, Cycles: 4,498,791, Regs: 32, GPU: TITAN V, SM Frequency: 1.19 cycle/nsecond, CC: 7.0, and Process: [925] 1-1-pinned-basic. Below these, there are two panes: "Instructions Executed" on the left showing source code from "1_1_pinned_basic.cu" and "Instructions Executed" on the right showing assembly instructions from "_Z6mygemmPfPKfS1_iii". A red arrow points from the text below to the "Resolve" button in the top toolbar between the two panes.

Hit “resolve” and load the corresponding source file from local file system

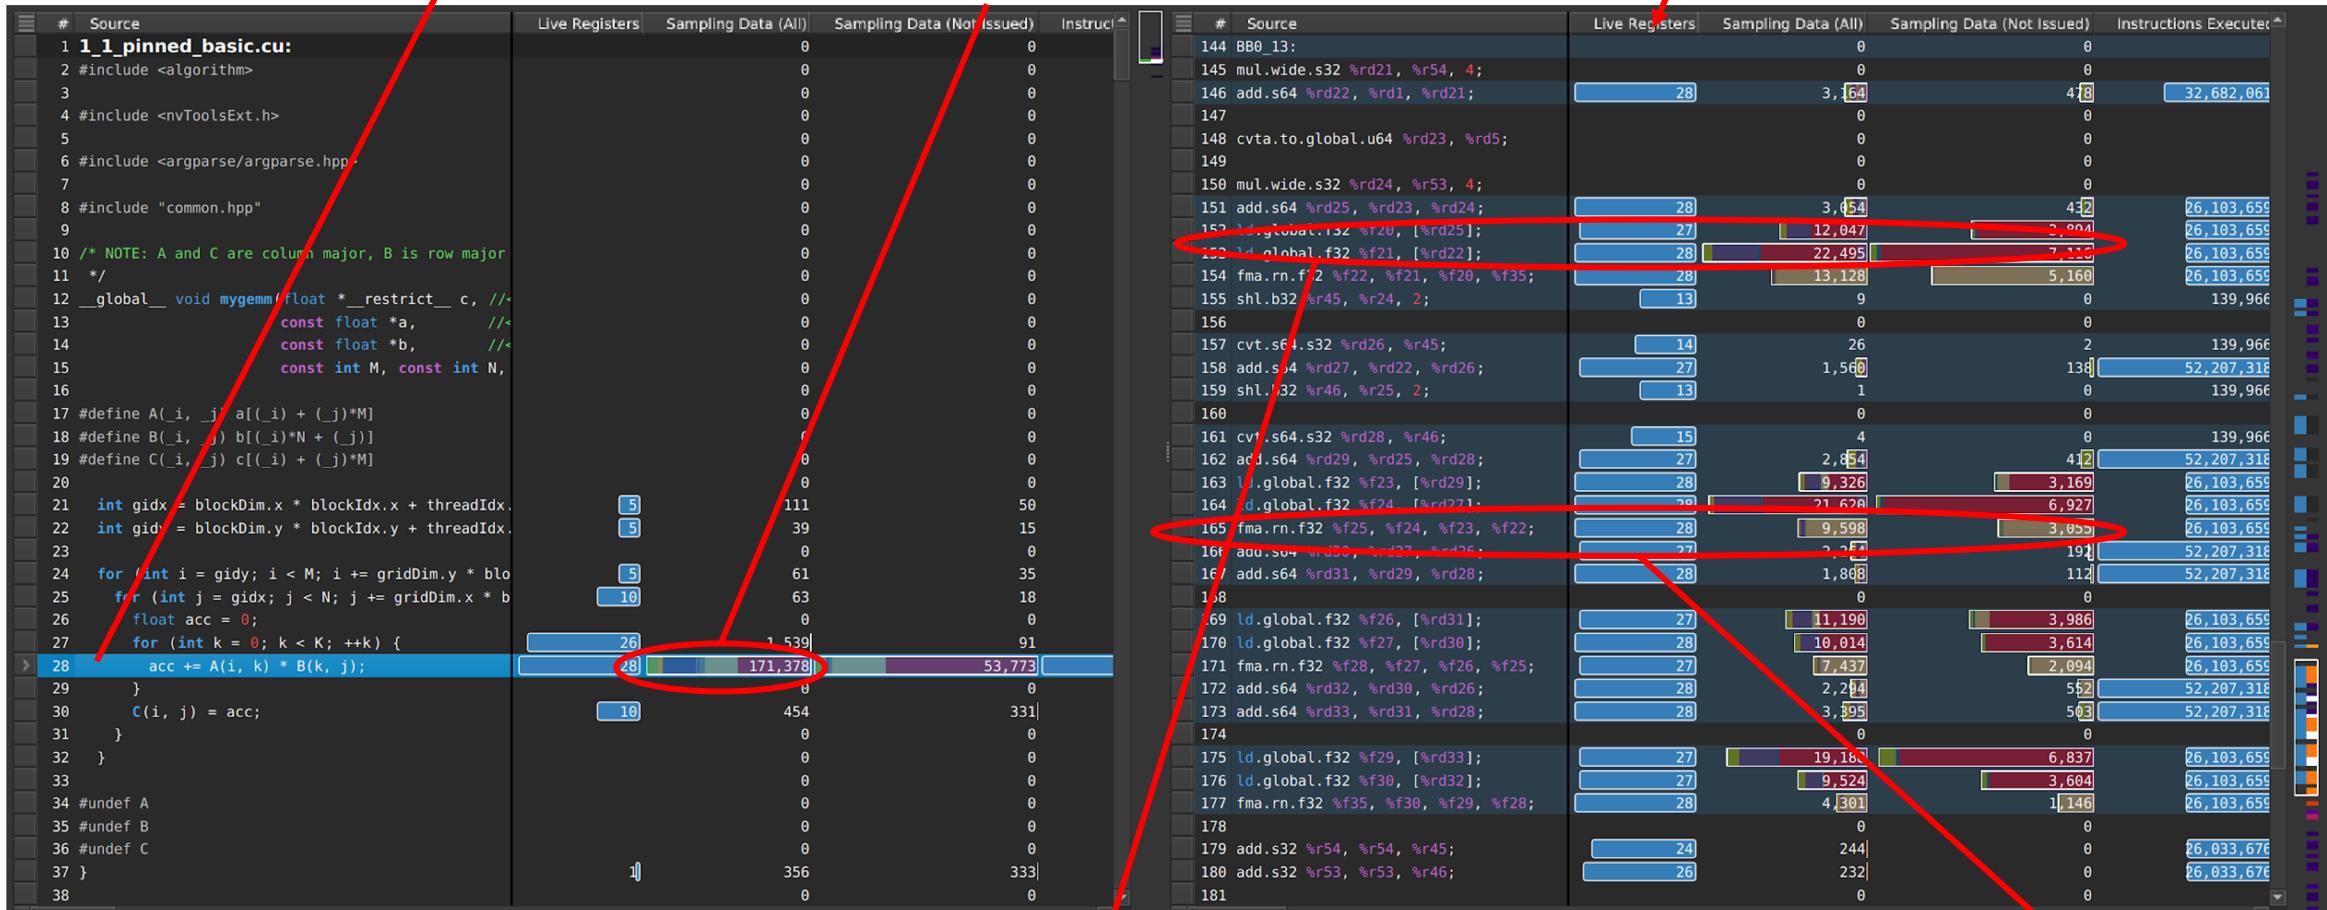
Instruction Sampling

- Every so often, the position of the program counter is recorded
- Slower instructions are more likely to be recorded
- There will be many samples in slow parts of the code, and few in fast parts of the code

Click a line to highlight lines from other side

Program counter spends most of its time on instructions from this line.
Mouse over for breakdown.

Corresponding PTX/SASS lines over here.



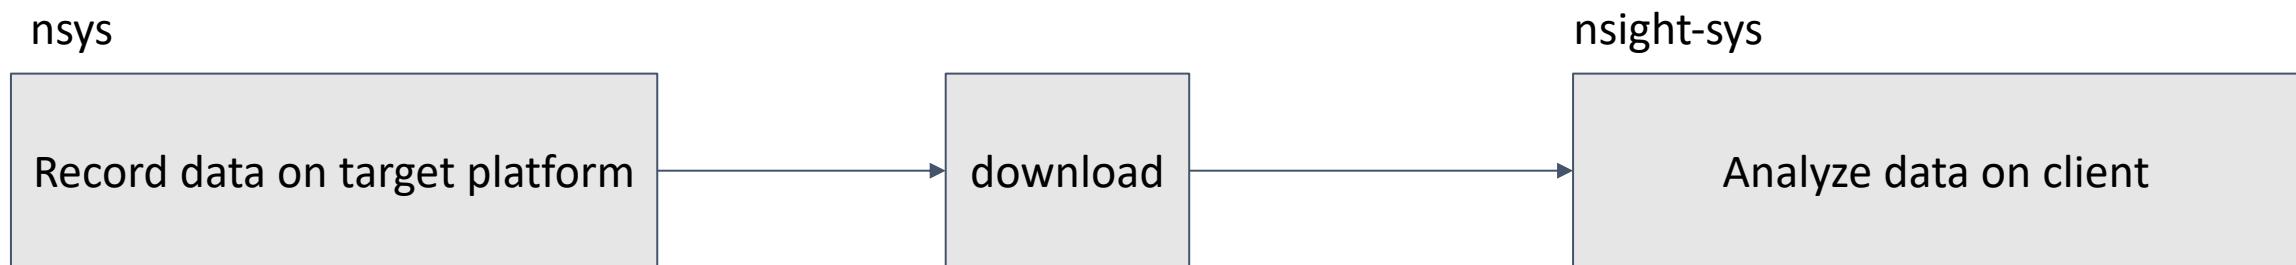
Multiple PTX lines per source line

Sometimes, stalls can show in a following instruction that depends on a previous one

System Profiling with Nsight Systems

Nvidia Nsight Systems

- Deliver work to the GPU effectively
 - Understand performance of surrounding system
- Two interfaces:
 - GUI (nsight-sys)
 - CLI (nsys)
- Like Nsight Compute, use a two-part record-then-analyze flow



Record Application Traces

```
$ nsys profile  
  -o 2-5-pinned-joint \\\n  2-5-pinned-joint
```

Create “2-5-pinned-basic.qdrep”
Name of the CUDA executable to profile

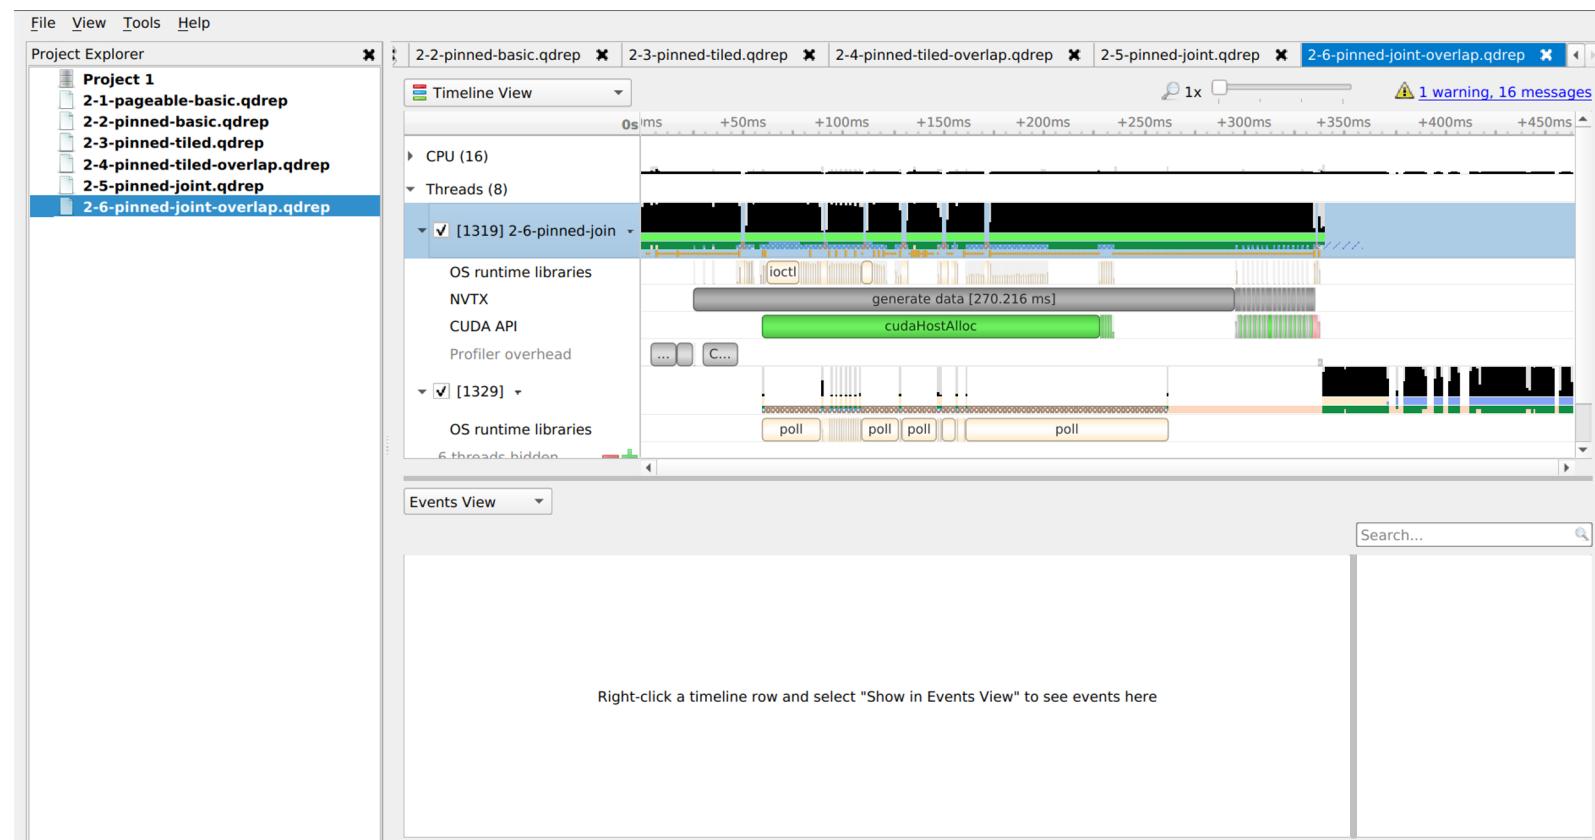
Nsight Systems GUI

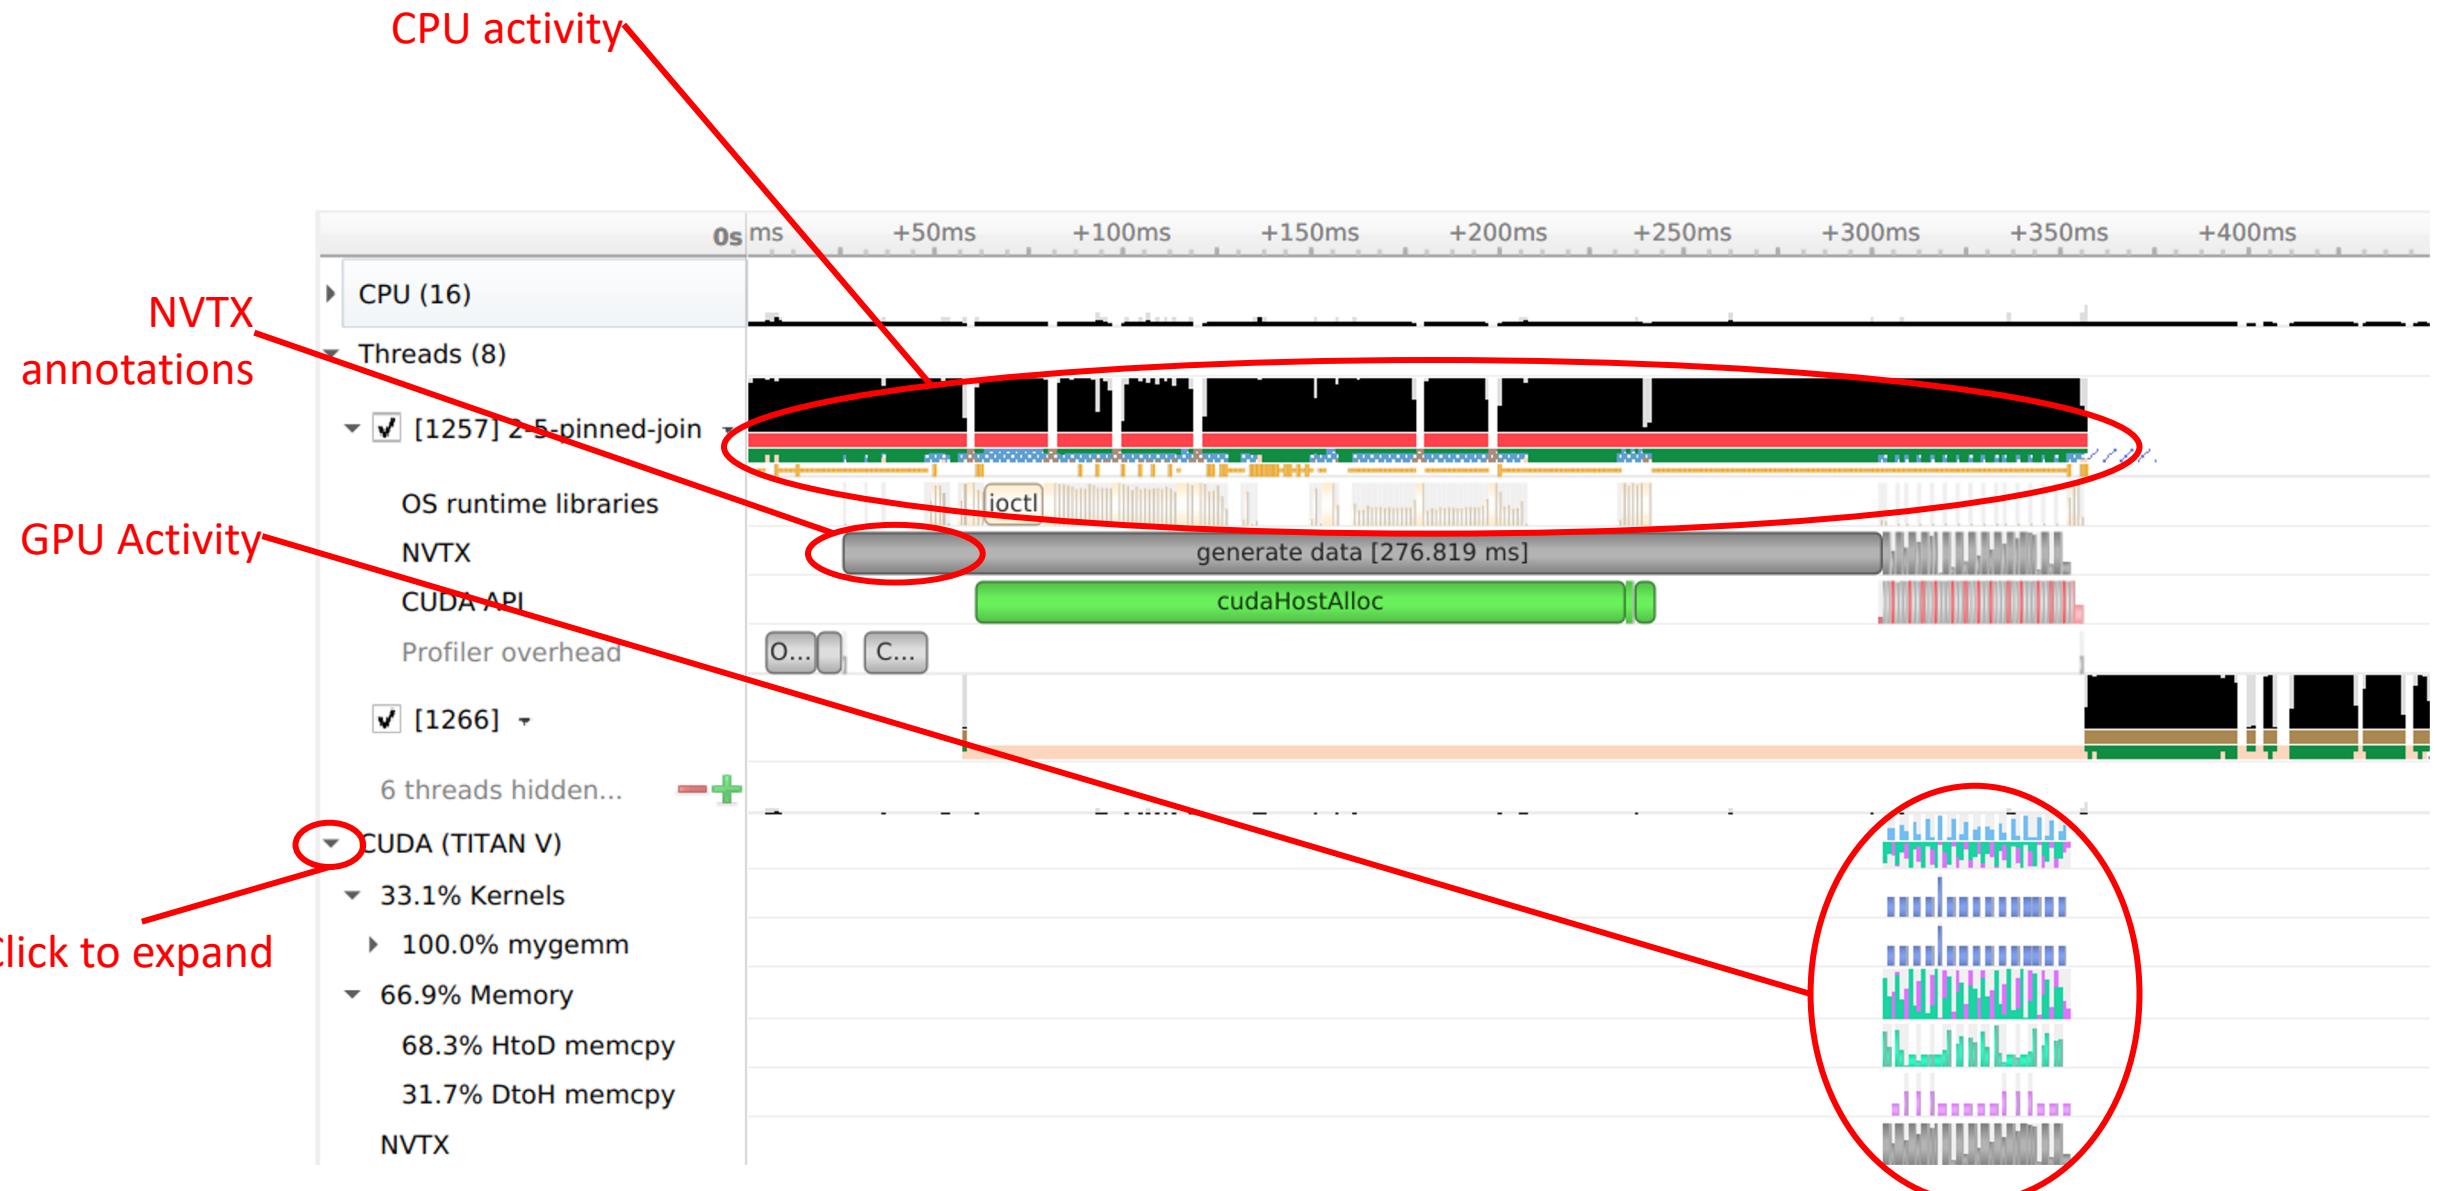
File > Open >
file.qdrqp

Multiple files will be open,
shown on the left pane.

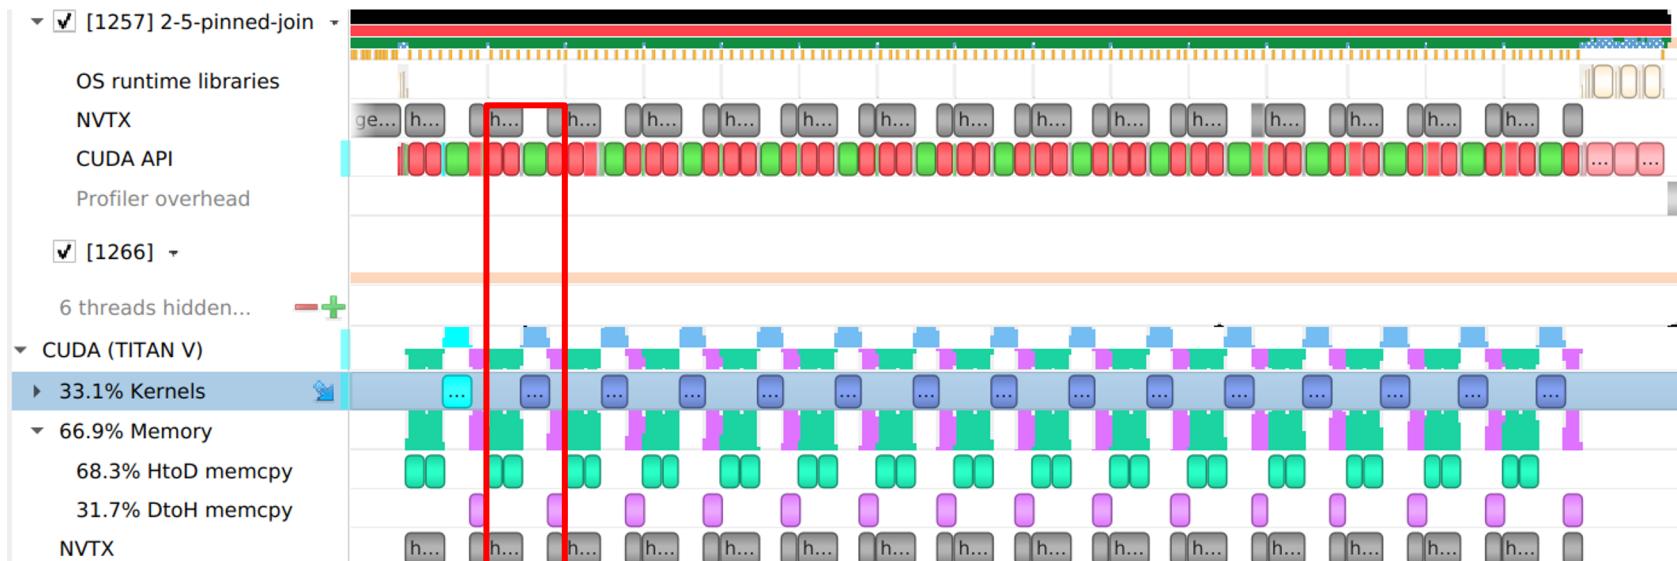
Main view is a timeline of
OS calls, CUDA calls, NVTX
events, CUDA API calls, and
GPU activity.

Open all the .qdrqp files
from the rai build directory
you downloaded.



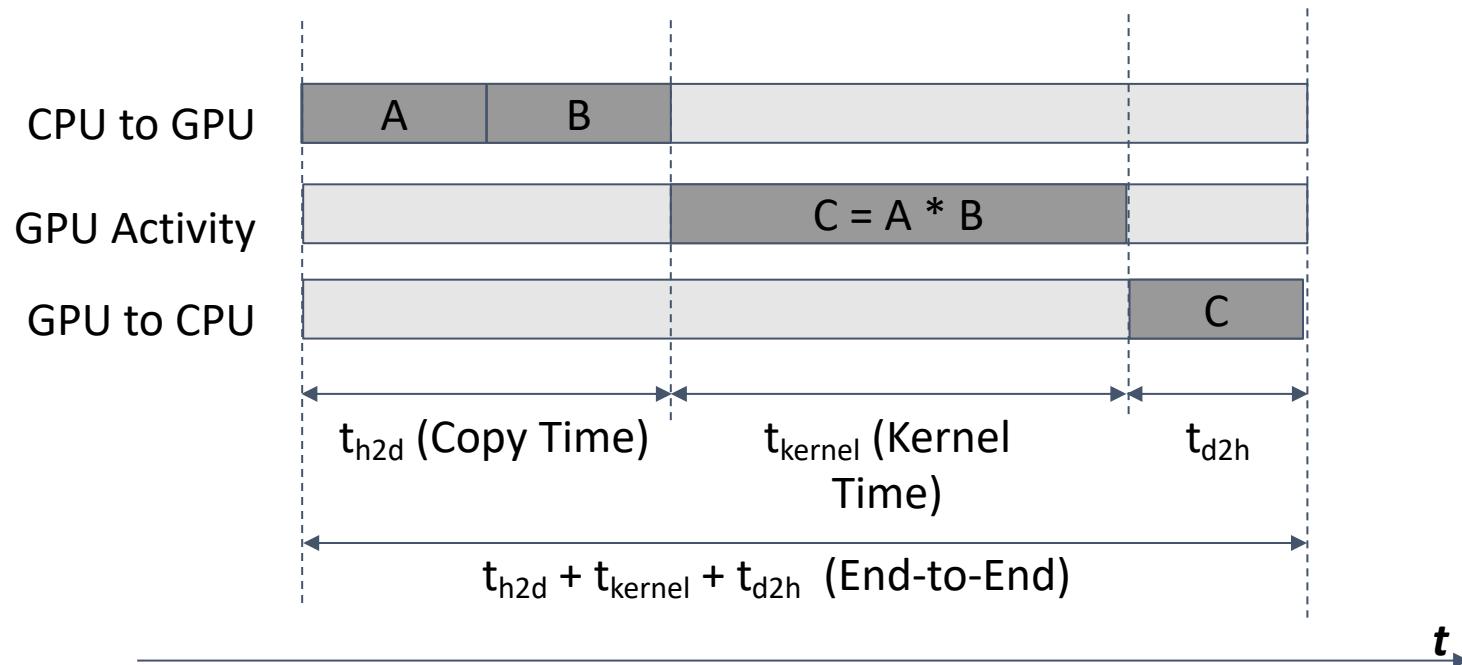


Click and drag to zoom in

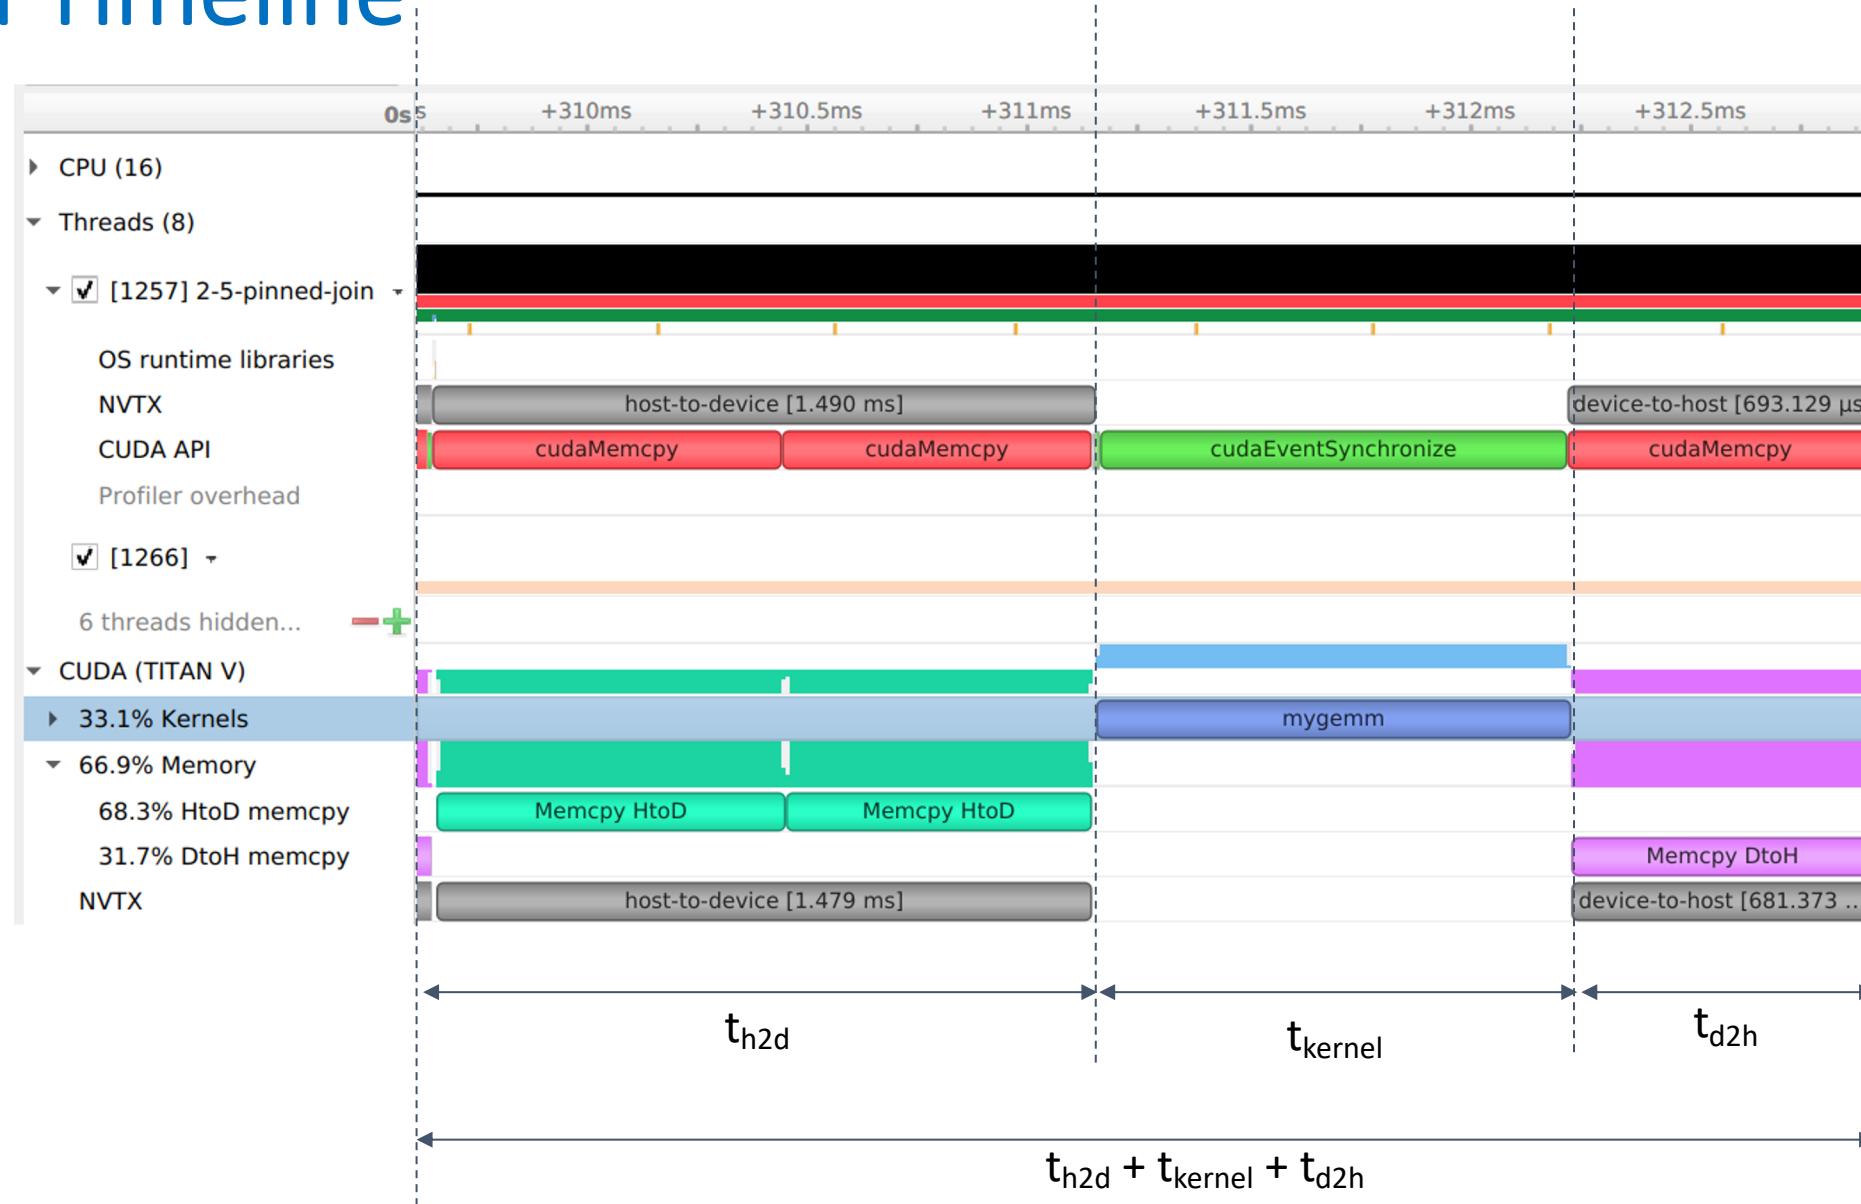


Mouse over spans for more info

Kernel Time vs End-to-End Time



Read Timeline



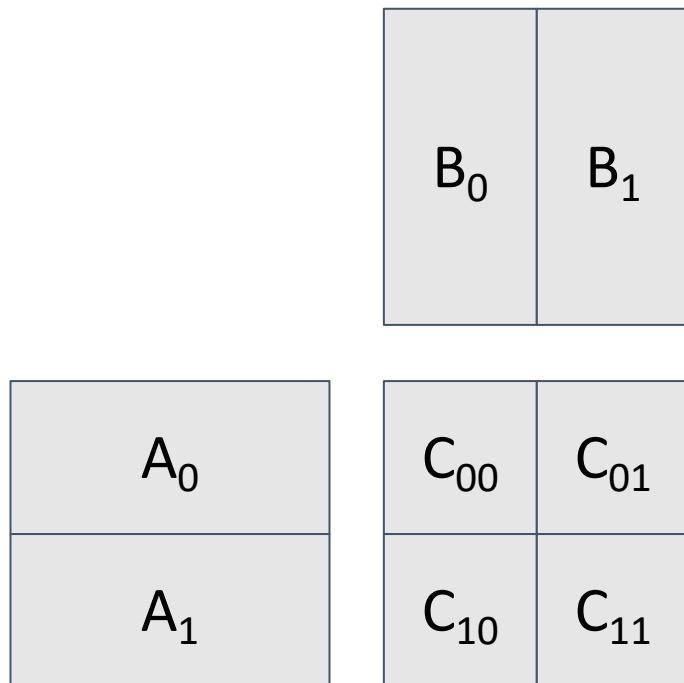
Overlap to reduce End-to-end Time

- $C = A * B$ as four multiplications.
 - $C_{00} = A_0 * B_0$: needs only A_0, B_0
 - $C_{01} = A_0 * B_1$: after C_{00} , needs only B_1
 - $C_{10} = A_1 * B_0$: after C_{01} , needs only A_1
 - $C_{11} = A_1 * B_1$: immediately after C_{10}
- Copy slices of A and B onto GPU, and immediately start the multiplication.
- Also can copy results back as soon as they're ready

A_0
A_1

C_{00}	
C_{01}	
C_{10}	C_{11}

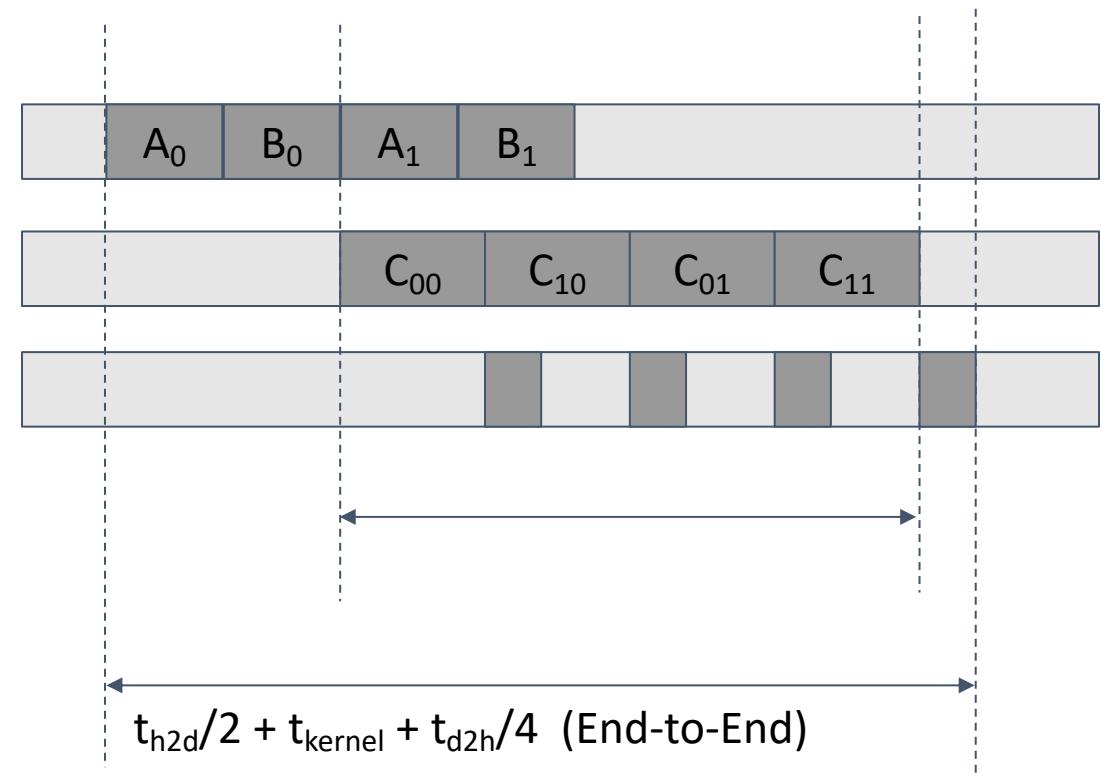
Overlap to Reduce End-to-End Time



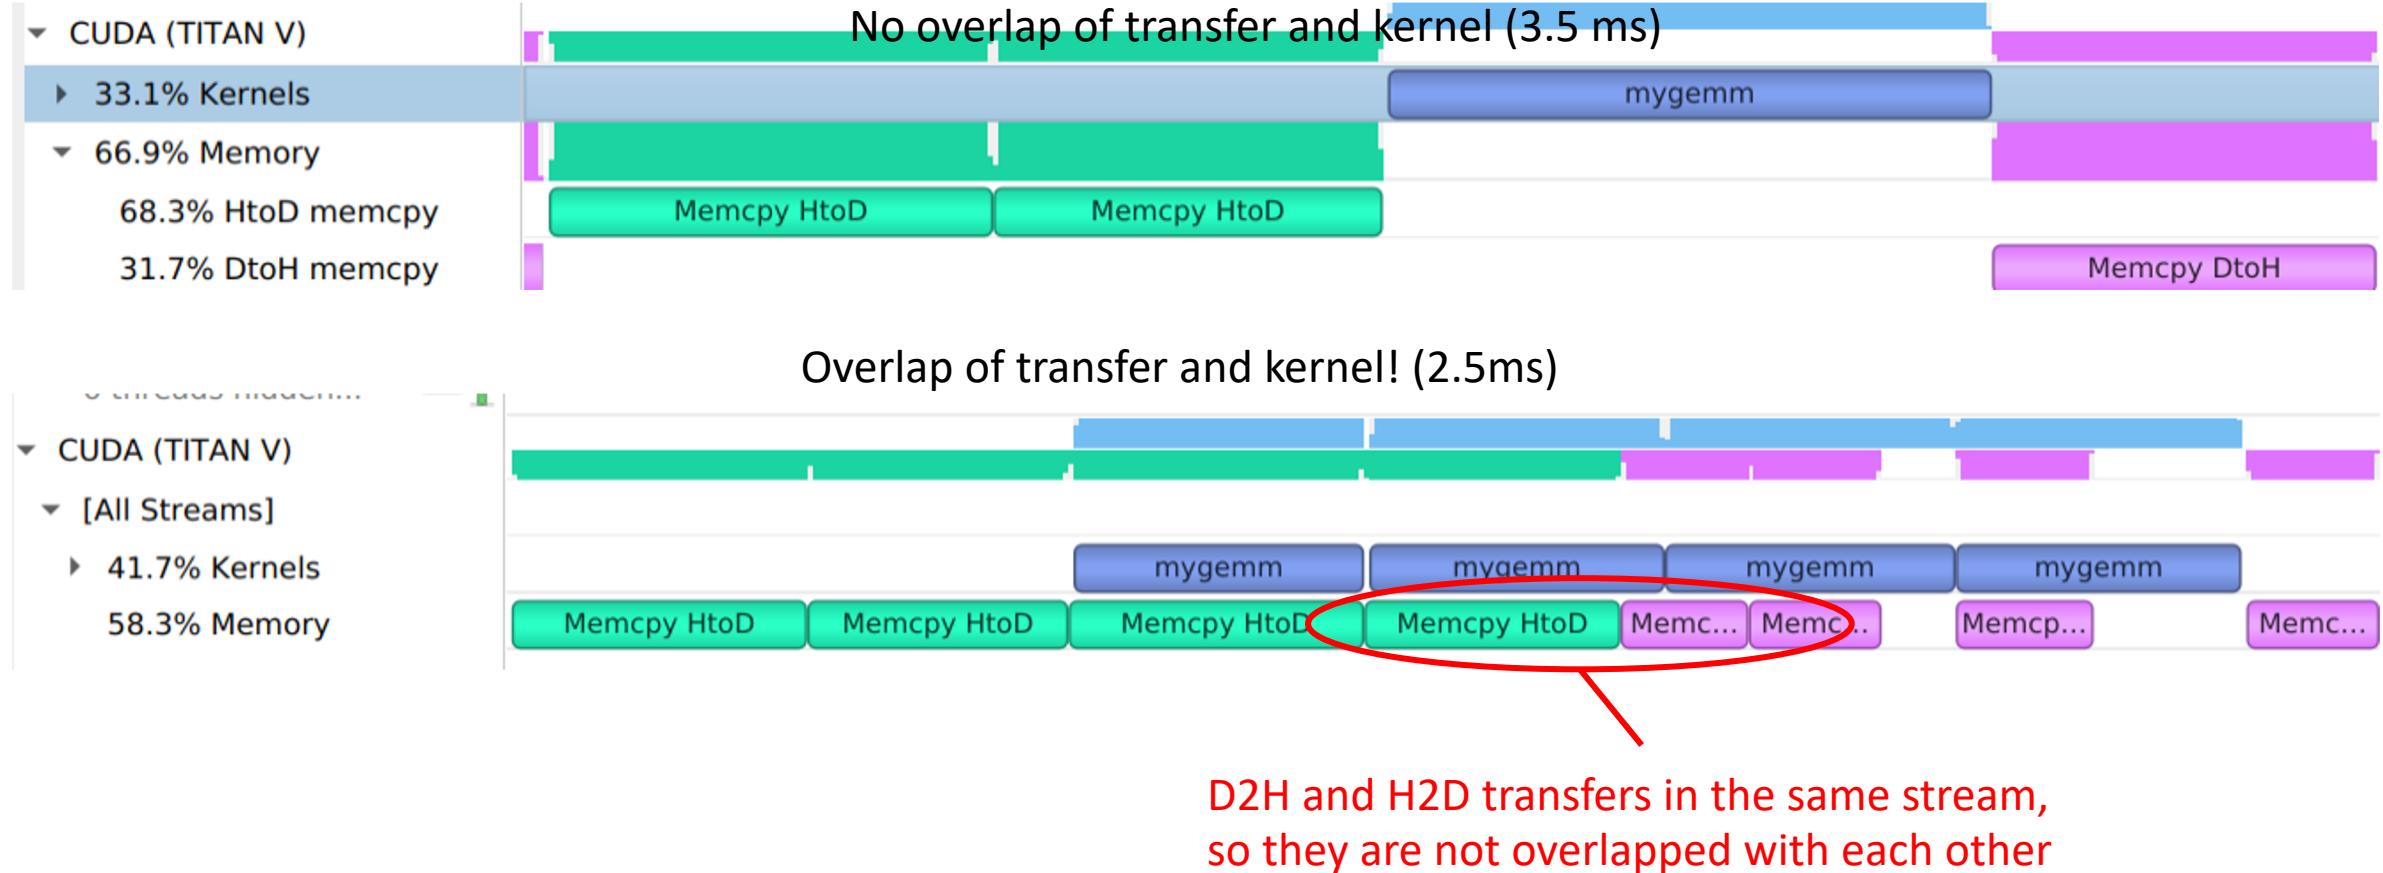
CPU to GPU

GPU Activity

GPU to CPU



Real Timelines: Overlapping



Further Reading

- Nsight Systems Documentation
 - <https://docs.nvidia.com/nsight-systems/>
- Nsight Compute Documentation
 - <https://docs.nvidia.com/nsight-compute/>
- Nvidia Developer Blog
 - Nsight Systems Exposes GPU Optimization (May 30 2018): <https://devblogs.nvidia.com/nsight-systems-exposes-gpu-optimization/>
 - Using Nsight Compute to Inspect your Kernels (Sep 16 2019): <https://devblogs.nvidia.com/using-nsight-compute-to-inspect-your-kernels/>
 - Using Nvidia Nsight Systems in Containers and the Cloud (Jan 29 2020) : <https://devblogs.nvidia.com/nvidia-nsight-systems-containers-cloud/>
- Workload Memory Analysis
 - CUDA Memory Model: <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#memory-hierarchy>
 - Device Memory Access Performance Guidelines: <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#device-memory-accesses>
- Stall Reasons
 - Nsight Graphics Docs: Stall Reasons: https://docs.nvidia.com/drive/drive_os_5.1.12.0L/nsight-graphics/activities/#shaderprofiler_stallreasons
 - Issue Efficiency Nsight Visual Studio Edition:
<https://docs.nvidia.com/gameworks/content/developertools/desktop/analysis/report/cudaexperiments/kernellevel/issueefficiency.htm>
- Occupancy:
 - Nsight Visual Studio Edition: <https://docs.nvidia.com/gameworks/content/developertools/desktop/analysis/report/cudaexperiments/kernellevel/achievedoccupancy.htm>

Not Discussed

- Measuring across multiple streams with CUDA events
- Profiling through the Nsight Compute GUI
 - local/remote
- Profiling through the Nsight Systems GUI
 - local/remote
- In-kernel timing with `clock()`/`clock64()`
- Custom profiling hooks with CUDA Performance Tools Interface (CUPTI)

Thank You

Copyright 2020 Carl Pearson

Material in these slides is just my own - Nvidia has not reviewed it.

The content of these slides may be reused or modified freely with attribution.